

Workshop Multitenant, Multimodel, In-Memory para la base de Datos Oracle

Parte 3 de 3



Contenidos

WORKSHOP MULTITENANT, MULTIMODEL, IN-MEMORY PARA LA BASE DE DATOS ORACLE	1
PARTE 3 DE 3.....	1
IN-MEMORY (30 MIN).....	3
CONFIGURACIÓN DEL ÁREA DE MEMORIA (5 MIN).....	3
<i>Configurar FastStart (5 min).....</i>	<i>6</i>
<i>Para desactivar el FAST START.....</i>	<i>7</i>
<i>Publicar las tablas SSB en in-memory (5 min).....</i>	<i>8</i>
<i>Monitorizar la publicación de SSB en In Memory.....</i>	<i>10</i>
<i>Queries Sencillas.....</i>	<i>11</i>
<i>Queries de grado medio.....</i>	<i>20</i>
<i>Queries complejas.....</i>	<i>28</i>
ACO - ORACLE ADVANCED COMPRESSION (45 MIN)	40
ADVANCED ROW COMPRESSION IMPLEMENTATION.....	40
ADVANCED ROW COMPRESSION TABLA/PARTICIÓN	45
<i>Utilización de API PL/SQL DBMS_COMRESSION.....</i>	<i>50</i>
<i>Proceso de compresión de índice.....</i>	<i>55</i>
ADVANCED COMPRESSION: ORACLE SECUREFILES	57
<i>Consultas sobre tablas Comprimidas/No Comprimidas</i>	<i>67</i>
<i>Proceso de compresión de expdp.....</i>	<i>69</i>



In-Memory (30 min)

Configuración del área de memoria (5 min)

Todas las prácticas de hoy se realizarán con la base de datos CDBA, por lo que debemos cargar sus variables de entorno al comienzo de las prácticas. Antes de comenzar, debemos asegurar que la instancia CDBB está parada. Además, dejaremos cargadas las variables de entorno de la instancia CDBA, dado que es la instancia donde realizaremos las prácticas hoy:

```
[oracle@single19c]$ . CDBB.env
[oracle@single19c]$ srvctl stop database -d CDBB -o immediate
[oracle@single19c]$ . CDBA.env
```

En primer lugar comprobamos la configuración de In-Memory que hay en la base de datos.

```
*****
A. In-Memory Column Store (IM column store) dynamic resizing:
*****
```

```
sqlplus / as sysdba
```

```
SQL> show parameter inmemo
```

NAME	TYPE	VALUE
inmemory_adg_enabled	boolean	TRUE
inmemory_automatic_level	string	OFF
inmemory_clause_default	string	
inmemory_expressions_usage	string	ENABLE
inmemory_force	string	DEFAULT
inmemory_max_populate_servers	integer	0
inmemory_optimized_arithmetic	string	DISABLE
inmemory_prefer_xmem_memcompress	string	
inmemory_prefer_xmem_priority	string	
inmemory_query	string	ENABLE
inmemory_size	big integer	0
inmemory_trickle_repopulate_servers_	integer	1
percent		
inmemory_virtual_columns	string	MANUAL
inmemory_xmem_size	big integer	0
optimizer_inmemory_aware	boolean	TRUE



```
SQL> select version from v$instance;
```

```
VERSION
-----
19.0.0.0.0
```

Para activar IMC, hay que poner `inmemory_size > 0` y reiniciar la instancia.
Aprovecharemos también para incrementar la sga de la misma:

```
SQL> show parameter sga
```

NAME	TYPE	VALUE
allow_group_access_to_sga	boolean	FALSE
lock_sga	boolean	FALSE
pre_page_sga	boolean	TRUE
sga_max_size	big integer	4G
sga_min_size	big integer	0
sga_target	big integer	4G
unified_audit_sga_queue_size	integer	1048576

```
SQL> alter system set sga_max_size=6G scope=spfile;
```

```
System altered.
```

```
SQL> alter system set sga_target=6G scope=spfile;
```

```
System altered.
```

```
SQL> alter system set inmemory_size = 2G scope=spfile;
```

```
System altered.
```

```
SQL> exit
```

```
[oracle@single19c ~]$ srvctl stop database -d $ORACLE_UNQNAME -o immediate
```

```
[oracle@single19c ~]$ srvctl start database -d $ORACLE_UNQNAME
```

```
$ sqlplus / as sysdba
```

```
SQL*Plus: Release 19.0.0.0.0 - Production on Wed Jan 12 11:13:37 2022  
Version 19.13.0.0.0
```

```
Copyright (c) 1982, 2021, Oracle. All rights reserved.
```

```
Connected to:
```

```
Oracle Database 19c EE Extreme Perf Release 19.0.0.0.0 - Production
```



Version 19.13.0.0.0

SQL> show parameter inmemo

NAME	TYPE	VALUE
inmemory_adg_enabled	boolean	TRUE
inmemory_automatic_level	string	OFF
inmemory_clause_default	string	
inmemory_expressions_usage	string	ENABLE
inmemory_force	string	DEFAULT
inmemory_max_populate_servers	integer	2
inmemory_optimized_arithmetic	string	DISABLE
inmemory_prefer_xmem_memcompress	string	
inmemory_prefer_xmem_priority	string	
inmemory_query	string	ENABLE
inmemory_size	big integer	2G
inmemory_trickle_repopulate_servers_percent	integer	1
inmemory_virtual_columns	string	MANUAL
inmemory_xmem_size	big integer	0
optimizer_inmemory_aware	boolean	TRUE

El resize puede ser dinámico :

SQL> alter system set inmemory_size = 3G scope=both;

System altered.

SQL> show parameter inmemory

NAME	TYPE	VALUE
inmemory_adg_enabled	boolean	TRUE
inmemory_automatic_level	string	OFF
inmemory_clause_default	string	
inmemory_expressions_usage	string	ENABLE
inmemory_force	string	DEFAULT
inmemory_max_populate_servers	integer	2
inmemory_optimized_arithmetic	string	DISABLE
inmemory_prefer_xmem_memcompress	string	
inmemory_prefer_xmem_priority	string	
inmemory_query	string	ENABLE
inmemory_size	big integer	3G
inmemory_trickle_repopulate_servers_percent	integer	1
inmemory_virtual_columns	string	MANUAL
inmemory_xmem_size	big integer	0
optimizer_inmemory_aware	boolean	TRUE

El proceso de resize dinámico se puede hacer solo al alza:



```
SQL> alter system set inmemory_size = 1G scope=both;
alter system set inmemory_size = 1G scope=both
*
ERROR at line 1:
ORA-02097: parameter cannot be modified because specified value is invalid
ORA-02095: specified initialization parameter cannot be modified
```

Configurar FastStart (5 min)

El área FastStart es un espacio de tablas designado donde IM FastStart almacena y gestiona los datos de los objetos INMEMORY. Oracle Database gestiona los Espacios de tablas FastStart automáticamente.

En una base de datos Oracle RAC, todos los nodos comparten los datos de FastStart.

```
[oracle@single19]$ sqlplus / as sysdba

SQL*Plus: Release 19.0.0.0.0 - Production on Thu Jan 13 09:04:06 2022
Version 19.13.0.0.0

Copyright (c) 1982, 2021, Oracle. All rights reserved.

Connected to:
Oracle Database 19c EE Extreme Perf Release 19.0.0.0.0 - Production
Version 19.13.0.0.0

SQL>
Connected.

col tablespace_name format a30

select con_id, TABLESPACE_NAME, STATUS FROM V$INMEMORY_FASTSTART_AREA;

   CON_ID TABLESPACE_NAME          STATUS
-----
1 INVALID_TABLESPACE              DISABLE
2 INVALID_TABLESPACE              DISABLE
3 INVALID_TABLESPACE              DISABLE
4 INVALID_TABLESPACE              DISABLE
5 INVALID_TABLESPACE              DISABLE

SQL> conn system/we1c0m3_we1c0m3_@SOE

Connected.

SQL> create tablespace TBS_IMC_FASTSTART datafile size 8G;

Tablespace created.

SQL> EXEC DBMS_INMEMORY_ADMIN.FASTSTART_ENABLE('TBS_IMC_FASTSTART')
```



PL/SQL procedure successfully completed.

SQL> conn / as sysdba

Connected.

col tablespace_name format a30

SQL> select con_id, TABLESPACE_NAME, STATUS FROM V\$INMEMORY_FASTSTART_AREA;

CON_ID	TABLESPACE_NAME	STATUS
1	INVALID_TABLESPACE	DISABLE
2	INVALID_TABLESPACE	DISABLE
3	INVALID_TABLESPACE	DISABLE
4	INVALID_TABLESPACE	DISABLE
5	TBS_IMC_FASTSTART	ENABLE

SQL> conn system/we1c0m3_we1c0m3_@SOE

Connected.

SQL> COL TABLESPACE_NAME FORMAT a20

SQL> SELECT TABLESPACE_NAME, STATUS,
((ALLOCATED_SIZE/1024) / 1024) AS ALLOC_MB,
((USED_SIZE/1024) / 1024) AS USED_MB
FROM V\$INMEMORY_FASTSTART_AREA;

TABLESPACE_NAME	STATUS	ALLOC_MB	USED_MB
TBS_IMC_FASTSTART	ENABLE	8192	1

Algunas notas sobre Fast Start:

- No se puede forzar de forma manual una escritura al FS !!!
- Se puede migrar el contenido del FS a otro TBS:

EXEC DBMS_INMEMORY_ADMIN.FASTSTART_MIGRATE_STORAGE('new_fs_tbs')

- Se puede deshabilitar el FS fastStart:

EXEC DBMS_INMEMORY_ADMIN.FASTSTART_DISABLE

Para desactivar el FAST START

SQL> conn system/we1c0m3_we1c0m3_@SOE

Connected.



```
SQL> EXEC DBMS_INMEMORY_ADMIN.FASTSTART_DISABLE
```

PL/SQL procedure successfully completed.

```
SQL> drop tablespace TBS_IMC_FASTSTART including contents and datafiles;
```

Tablespace dropped.

Publicar las tablas SSB en in-memory (5 min)

```
SQL> conn ssb/ssb@SOE
```

Connected.

```
col table_name format a30  
set lines 120
```

```
--display current status
```

```
SQL> select table_name,  
            inmemory,  
            inmemory_priority,  
            inmemory_compression  
from user_tables;
```

TABLE_NAME	INMEMORY	INMEMORY	INMEMORY_COMPRESS
LINEORDER	DISABLED		
LINEORDER_ACO	DISABLED		
ETL_LO	DISABLED		
RESULTS	DISABLED		
DATE_DIM	DISABLED		
ETL_DD	DISABLED		
SUPPLIER	DISABLED		
PART	DISABLED		
CUSTOMER	DISABLED		
TMP	DISABLED		
YEARLY_PROFIT_REP_MV		DISABLED	
LINEORDER_NO_ACO	DISABLED		

12 rows selected.

```
--alter tables in memory
```

```
SQL> alter table lineorder inmemory;  
alter table part inmemory;  
alter table customer inmemory;  
alter table supplier inmemory;  
alter table date_dim inmemory;
```

Table altered.



Table altered.

Table altered.

Table altered.

Table altered.

```
SQL> select table_name,
        inmemory,
        inmemory_priority,
        inmemory_compression
from user_tables;
```

TABLE_NAME	INMEMORY	INMEMORY	INMEMORY_COMPRESS
LINEORDER	ENABLED	NONE	FOR QUERY LOW
LINEORDER_ACO	DISABLED		
ETL_LO	DISABLED		
RESULTS	DISABLED		
DATE_DIM	ENABLED	NONE	FOR QUERY LOW
ETL_DD	DISABLED		
SUPPLIER	ENABLED	NONE	FOR QUERY LOW
PART	ENABLED	NONE	FOR QUERY LOW
CUSTOMER	ENABLED	NONE	FOR QUERY LOW
TMP	DISABLED		
YEARLY_PROFIT_REP_MV		DISABLED	
LINEORDER_NO_ACO	DISABLED		

11 rows selected.

--fetch all rows to start population

```
SQL> select count(*) from lineorder;
select count(*) from part;
select count(*) from customer;
select count(*) from supplier;
select count(*) from date_dim;
```

```
    COUNT(*)
-----
    29999870
```

```
    COUNT(*)
-----
    600000
```

```
    COUNT(*)
-----
    150000
```

```
    COUNT(*)
-----
    10000
```



```
COUNT(*)
```

```
-----  
2560
```

Monitorizar la publicación de SSB en In Memory

```
SQL> conn system/we1c0m3_we1c0m3_@SOE
```

```
Connected.
```

```
--new view v$im_segments
```

```
SQL> desc v$im_segments
```

Name	Null?	Type
OWNER		VARCHAR2(128)
SEGMENT_NAME		VARCHAR2(128)
PARTITION_NAME		VARCHAR2(128)
SEGMENT_TYPE		VARCHAR2(18)
TABLESPACE_NAME		VARCHAR2(128)
INMEMORY_SIZE		NUMBER
BYTES		NUMBER
BYTES_NOT_POPULATED		NUMBER
POPULATE_STATUS		VARCHAR2(13)
INMEMORY_PRIORITY		VARCHAR2(8)
INMEMORY_DISTRIBUTE		VARCHAR2(15)
INMEMORY_DUPLICATE		VARCHAR2(13)
INMEMORY_COMPRESSION		VARCHAR2(17)
INMEMORY_SERVICE		VARCHAR2(12)
INMEMORY_SERVICE_NAME		VARCHAR2(129)
IS_EXTERNAL		VARCHAR2(5)
CON_ID		NUMBER

```
col owner format a12
```

```
col name format a30
```

```
col partition_name format a30
```

```
set lines 120
```

```
--population status
```

```
SQL> select v.owner,v.segment_name name,v.partition_name,  
v.populate_status status, v.bytes_not_populated  
from v$im_segments v  
Order by 1;
```

OWNER	NAME	PARTITION_NAME	STATUS
BYTES_NOT_POPULATED			



```

-----
SSB      LINEORDER      STARTED
      401473536
SSB      PART      COMPLETED
      0
SSB      CUSTOMER      COMPLETED
      0
SSB      SUPPLIER      COMPLETED
      0
SSB      DATE_DIM      COMPLETED
      0

```

6 rows selected.

(La anterior query se puede ejecutar varias veces para ver como van quedando menos bytes por subir a memoria)

--size

```

SQL> select v.owner,v.segment_name name,
round(v.bytes/1024/1024,3) orig_size,
round(v.inmemory_size/1024/1024,3) in_mem_size,
ROUND(v.bytes/v.inmemory_size,2) comp_ratio
from v$im_segments v
order by 1;

```

OWNER	NAME	ORIG_SIZE	IN_MEM_SIZE	COMP_RATIO	
SSB	SUPPLIER	.836	1.25	.67	
SSB	DATE_DIM	.117	1.25	.09	
SSB	LINEORDER		1745.32	1494.375	1.17
SSB	CUSTOMER	11.852	11.25	1.05	
SSB	PART	40.563	13.438	3.02	

6 rows selected.

Queries Sencillas

Comprobar la diferencia de acceso entre las distintas queries.

Single Table Scan

Query 1:

```

SQL> conn ssb/ssb@SOE

```

Connected.

```

/*****Query 1*****/
/*****Single Column Aggregation*****/
/*****

```



```

/*****Parallel Disk Access*****/
/*****/

clear scr

--flush the buffer_cache

SQL> alter system flush buffer_cache;
SQL> alter session force parallel query parallel 4;

System altered.

Session altered.

set lines 120
set autotrace traceonly explain statistics
set timing on

select /*+ NO_INMEMORY */ /* DISK ACCESS */
max(lo_ordtotalprice) most_expensive_order From LINEORDER;

```

Elapsed: 00:00:32.22

Execution Plan

Plan hash value: 396151021

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
TQ	IN-OUT PQ Distrib					
0	SELECT STATEMENT		1	6	17006 (1)	
1	SORT AGGREGATE		1	6		
2	PX COORDINATOR					
3	PX SEND QC (RANDOM)	:TQ10000	1	6		
Q1,00	P->S QC (RAND)					
4	SORT AGGREGATE		1	6		
Q1,00	PCWP					
5	PX BLOCK ITERATOR		29M	171M	17006 (1)	
Q1,00	PCWC					
6	TABLE ACCESS FULL	LINEORDER	29M	171M	17006 (1)	
Q1,00	PCWP					

Note

- Degree of Parallelism is 4 because of session



Statistics

```

-----
1338 recursive calls
0 db block gets
225488 consistent gets
223484 physical reads
5548 redo size
565 bytes sent via SQL*Net to client
463 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
110 sorts (memory)
0 sorts (disk)
1 rows processed

```

```

/*****In-Memory Serial Access*****/
/*****/

SQL> clear scr

--column store enabled via INMEMORY_QUERY parameter--

SQL> alter session disable parallel query;

Session altered

SQL> Select /*+ INMEMORY */ /*IN-MEMORY Serial*/
max(lo_ordtotalprice) most_expensive_order From LINEORDER;

Elapsed: 00:00:00.15

Execution Plan
-----
Plan hash value: 2267213921

-----
| Id | Operation | Name | Rows | Bytes | Cost (%CPU)|
Time |          |      |      |      |             |
-----
| 0 | SELECT STATEMENT |      | 1 | 6 | 2561 (12)|
00:00:01 |
| 1 | SORT AGGREGATE |      | 1 | 6 |           |
|
| 2 | TABLE ACCESS INMEMORY FULL | LINEORDER | 29M | 171M | 2561 (12)|
00:00:01 |
-----

```



Statistics

```

-----
1252 recursive calls
   8 db block gets
2156 consistent gets
   4 physical reads
1408 redo size
565 bytes sent via SQL*Net to client
463 bytes received via SQL*Net from client
   2 SQL*Net roundtrips to/from client
125 sorts (memory)
   0 sorts (disk)
   1 rows processed

```

```

/*****In-Memory Parallel Access*****/
/*****/

SQL> clear scr

SQL> alter session force parallel query parallel 4;

Session altered.

SQL> select /*+ INMEMORY */ /*IN-MEMORY Parallel*/
max(lo_ordtotalprice) most_expensive_order From LINEORDER;

```

Elapsed: 00:00:00.34

Execution Plan

Plan hash value: 396151021

```

-----
--

```

Id	Operation	Name	Rows	Bytes	Cost
(%CPU)	Time	TQ	IN-OUT	PQ Distrib	
0	SELECT STATEMENT		1	6	711
(12)	00:00:01				
1	SORT AGGREGATE		1	6	

```

-----
--

```



--

— — — — —

- ## Statistics

Query 2

```

/*****Parallel Disk Access*****/

```



```

/*****
SQL> clear scr
--flush the buffer_cache
SQL> alter system flush buffer_cache;
SQL> alter session force parallel query parallel 4;
SQL> select /*+ NO_INMEMORY */ /* DISK ACCESS */
max(lo_ordtotalprice) most_expensive_order From LINEORDER
where LO_PARTKEY=300023;
Elapsed: 00:00:30.27
Execution Plan
-----
Plan hash value: 396151021
-----
-----
| Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |
|----|-----|-----|-----|-----|-----|-----|
-----
| 0 | SELECT STATEMENT | | 1 | 11 | 16966 (1)| 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 11 | | |
| 2 | PX COORDINATOR | | | | | |
| 3 | PX SEND QC (RANDOM) | :TQ10000 | 1 | 11 | | |
| Q1,00 | P->S | QC (RAND) | | | | |
| 4 | SORT AGGREGATE | | 1 | 11 | | |
| Q1,00 | PCWP | | | | | |
| 5 | PX BLOCK ITERATOR | | 76 | 836 | 16966 (1)| 00:00:01 |
| Q1,00 | PCWC | | | | | |
|* 6 | TABLE ACCESS FULL | LINEORDER | 76 | 836 | 16966 (1)| 00:00:01 |
| Q1,00 | PCWP | | | | | |
-----
-----
Predicate Information (identified by operation id):
-----
6 - filter("LO_PARTKEY"=300023)
Note
-----
- Degree of Parallelism is 4 because of session
Statistics
-----

```




```

1275 recursive calls
    7 db block gets
225541 consistent gets
223510 physical reads
  8540 redo size
    565 bytes sent via SQL*Net to client
    487 bytes received via SQL*Net from client
      2 SQL*Net roundtrips to/from client
    121 sorts (memory)
      0 sorts (disk)
      1 rows processed

```

```

/*****In-Memory Serial Access*****/
/*****/

SQL> clear scr

SQL> alter session disable parallel query;

SQL> Select /*+ INMEMORY */ /*IN-MEMORY Serial*/
max(lo_ordtotalprice) most_expensive_order From LINEORDER
where LO_PARTKEY=300023;

```

Elapsed: 00:00:00.09

Execution Plan

Plan hash value: 2267213921

Id	Operation	Name	Rows	Bytes	Cost (%CPU)
0	SELECT STATEMENT		1	11	2814 (20)
1	SORT AGGREGATE		1	11	
* 2	TABLE ACCESS INMEMORY FULL	LINEORDER	76	836	2814 (20)

Predicate Information (identified by operation id):



```
2 - inmemory("LO_PARTKEY"=300023)
   filter("LO_PARTKEY"=300023)
```

Statistics

```
-----
736 recursive calls
  0 db block gets
827 consistent gets
  1 physical reads
120 redo size
565 bytes sent via SQL*Net to client
487 bytes received via SQL*Net from client
  2 SQL*Net roundtrips to/from client
 82 sorts (memory)
  0 sorts (disk)
  1 rows processed
```

```
/*****In-Memory Parallel Access*****/
/*****/
```

```
SQL> clear scr
```

```
SQL> alter session force parallel query parallel 4;
```

```
SQL> select /*+ INMEMORY */ /*IN-MEMORY Parallel*/
max(lo_ordtotalprice) most_expensive_order From LINEORDER
where LO_PARTKEY=300023;
```

```
Elapsed: 00:00:00.05
```

Execution Plan

```
-----
Plan hash value: 396151021
```

```
-----
--
--
| Id | Operation | Name | Rows | Bytes | Cost |
| (%CPU)| Time | TQ | IN-OUT | PQ Distrib |
|-----|-----|-----|-----|-----|-----|
--
| 0 | SELECT STATEMENT | | 1 | 11 | 782 |
(20)| 00:00:01 | | | |
|
```



1	SORT AGGREGATE		1	11
2	PX COORDINATOR			
3	PX SEND QC (RANDOM) Q1,00 P->S QC (RAND)	:TQ10000	1	11
4	SORT AGGREGATE Q1,00 PCWP		1	11
5	PX BLOCK ITERATOR (20) 00:00:01 Q1,00 PCWC		76	836 782
* 6	TABLE ACCESS INMEMORY FULL LINEORDER (20) 00:00:01 Q1,00 PCWP		76	836 782

--

Predicate Information (identified by operation id):

6 - inmemory("LO_PARTKEY"=300023)
filter("LO_PARTKEY"=300023)

Note

- Degree of Parallelism is 4 because of session

Statistics

13 recursive calls
0 db block gets
115 consistent gets
0 physical reads
0 redo size
565 bytes sent via SQL*Net to client
489 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
0 sorts (memory)
0 sorts (disk)
1 rows processed



Queries de grado medio

Two table scan

Query 1

```
SQL> conn ssb/ssb@SOE

/*****Group Aggregate with two table join*****/
/*****

/*****Parallel Disk Access*****/
/*****

SQL> clear scr

--flush the buffer_cache

SQL> alter system flush buffer_cache;
SQL> alter session force parallel query parallel 4;
SQL> alter session set inmemory_query='DISABLE';
SQL> set autotrace traceonly explain statistics
SQL> Set timing on

SQL> select /*+ NO_INMEMORY */ /* DISK ACCESS */
d_date,sum(l.lo_revenue) "Total Revenue"
From   LINEORDER l, DATE_DIM d
Where  l.lo_orderdate = d.d_datekey
and    D_DAYNUMINMONTH = 25
and    d.d_month = 'December'
group by d_date
order by d_date;

6 rows selected.

Elapsed: 00:00:30.12

Execution Plan
-----
Plan hash value: 803921888

-----
-----
-----

| Id | Operation | Name |
Rows | Bytes | Cost (%CPU)| Time |
TQ | IN-OUT| PQ Distrib |
-----
-----
```



	0		SELECT STATEMENT			
7		343	17032 (2) 00:00:01			
	1		TEMP TABLE TRANSFORMATION			
	2		LOAD AS SELECT (CURSOR DURATION MEMORY)		SYS_TEMP_0FD9D6651_290D25	
	3		PX COORDINATOR			
	4		PX SEND QC (RANDOM)		:TQ10001	
7		294	5 (20) 00:00:01			
Q1,01		P->S	QC (RAND)			
	5		HASH GROUP BY			
7		294	5 (20) 00:00:01			
Q1,01		PCWP				
	6		PX RECEIVE			
7		294	4 (0) 00:00:01			
Q1,01		PCWP				
	7		PX SEND HASH		:TQ10000	
7		294	4 (0) 00:00:01			
Q1,00		P->P	HASH			
	8		KEY VECTOR CREATE BUFFERED		:KV0000	
7		294	4 (0) 00:00:01			
Q1,00		PCWC				
	9		PX BLOCK ITERATOR			
7		266	4 (0) 00:00:01			
Q1,00		PCWC				
	* 10		TABLE ACCESS FULL		DATE_DIM	
7		266	4 (0) 00:00:01			
Q1,00		PCWP				
	11		PX COORDINATOR			
	12		PX SEND QC (ORDER)		:TQ20004	
7		343	17026 (2) 00:00:01			
Q2,04		P->S	QC (ORDER)			



13	SORT ORDER BY		
7 343 17026 (2) 00:00:01			
Q2,04 PCWP			
14	PX RECEIVE		
7 343 17025 (2) 00:00:01			
Q2,04 PCWP			
15	PX SEND RANGE	:TQ20003	
7 343 17025 (2) 00:00:01			
Q2,03 P->P RANGE			
* 16	HASH JOIN BUFFERED		
7 343 17025 (2) 00:00:01			
Q2,03 PCWP			
17	PX RECEIVE		
7 119 17023 (2) 00:00:01			
Q2,03 PCWP			
18	PX SEND HYBRID HASH	:TQ20001	
7 119 17023 (2) 00:00:01			
Q2,01 P->P HYBRID HASH			
19	STATISTICS COLLECTOR		
Q2,01 PCWC			
20	VIEW	VW_VT_80F21617	
7 119 17023 (2) 00:00:01			
Q2,01 PCWP			
21	HASH GROUP BY		
7 112 17023 (2) 00:00:01			
Q2,01 PCWP			
22	PX RECEIVE		
7 112 17023 (2) 00:00:01			
Q2,01 PCWP			
23	PX SEND HASH	:TQ20000	
7 112 17023 (2) 00:00:01			
Q2,00 P->P HASH			
24	VECTOR GROUP BY		
7 112 17023 (2) 00:00:01			
Q2,00 PCWP			
25	HASH GROUP BY		
7 112 17023 (2) 00:00:01			
Q2,00 PCWP			
26	KEY VECTOR USE	:KV0000	
87245 1363K 17023 (2) 00:00:01			



Q2,00	PCWC				
27			PX BLOCK ITERATOR		
29M	343M	17023	(2)	00:00:01	
Q2,00	PCWC				
* 28			TABLE ACCESS FULL	LINEORDER	
29M	343M	17023	(2)	00:00:01	
Q2,00	PCWP				
29			PX RECEIVE		
7	224	2	(0)	00:00:01	
Q2,03	PCWP				
30			PX SEND HYBRID HASH	:TQ20002	
7	224	2	(0)	00:00:01	
Q2,02	P->P		HYBRID HASH		
31			PX BLOCK ITERATOR		
7	224	2	(0)	00:00:01	
Q2,02	PCWC				
32			TABLE ACCESS FULL	SYS_TEMP_0FD9D6651_290D25	
7	224	2	(0)	00:00:01	
Q2,02	PCWP				

Predicate Information (identified by operation id):

10 - filter("D_DAYNUMINMONTH"=25 AND "D"."D_MONTH"='December')
16 - access("ITEM_5"=INTERNAL_FUNCTION("C0"))
28 - filter(SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000))

Note

-
- dynamic statistics used: dynamic sampling (level=5)
 - Degree of Parallelism is 4 because of session
 - vector transformation used for this statement

Statistics

1182	recursive calls
43	db block gets
225566	consistent gets
223559	physical reads
18976	redo size
860	bytes sent via SQL*Net to client
604	bytes received via SQL*Net from client



2	SQL*Net roundtrips to/from client
271	sorts (memory)
0	sorts (disk)
6	rows processed

/*****In-Memory Serial Access*****/

SQL> clear scr

SQL> alter session disable parallel query;

SQL> alter session set inmemory_query='ENABLE';

SQL> Select /*+ INMEMORY */ /*IN-MEMORY Serial*/
d_date,sum(l.lo_revenue) "Total Revenue"
From LINEORDER l, DATE_DIM d
Where l.lo_orderdate = d.d_datekey
and D_DAYNUMINMONTH = 25
and d.d_month = 'December'
group by d_date
order by d_date;

6 rows selected.

Elapsed: 00:00:00.08

Execution Plan

Plan hash value: 2937007760

Id		Operation			Name	
Rows		Bytes		Cost (%CPU)	Time	

		0		SELECT STATEMENT		
7		343		2724 (17)	00:00:01	
		1		TEMP TABLE TRANSFORMATION		
		2		LOAD AS SELECT (CURSOR DURATION MEMORY)		SYS_TEMP_0FD9D6655_290D25
		3		HASH GROUP BY		
7		294		2 (50)	00:00:01	
		4		KEY VECTOR CREATE BUFFERED		:KV0000
7		294		1 (0)	00:00:01	
*		5		TABLE ACCESS INMEMORY FULL		DATE_DIM
7		266		1 (0)	00:00:01	



6		SORT GROUP BY	
7	343	2722 (17)	00:00:01
* 7		HASH JOIN	
7	343	2721 (17)	00:00:01
8		VIEW	VW_VT_80F21617
7	119	2719 (17)	00:00:01
9		VECTOR GROUP BY	
7	112	2719 (17)	00:00:01
10		HASH GROUP BY	
7	112	2719 (17)	00:00:01
11		KEY VECTOR USE	:KV0000
87245	1363K	2719 (17)	00:00:01
* 12		TABLE ACCESS INMEMORY FULL	LINEORDER
29M	343M	2719 (17)	00:00:01
13		TABLE ACCESS FULL	SYS_TEMP_0FD9D6655_290D25
7	224	2 (0)	00:00:01

Predicate Information (identified by operation id):

- 5 - inmemory("D_DAYNUMINMONTH"=25 AND "D"."D_MONTH"='December')
- filter("D_DAYNUMINMONTH"=25 AND "D"."D_MONTH"='December')
- 7 - access("ITEM_5"=INTERNAL_FUNCTION("C0"))
- 12 - inmemory(SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000))
- filter(SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000))

Note

- vector transformation used for this statement

Statistics

```

12 recursive calls
5 db block gets
20 consistent gets
0 physical reads
1096 redo size
860 bytes sent via SQL*Net to client
604 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
2 sorts (memory)
0 sorts (disk)
6 rows processed

```

```

/*****In-Memory Parallel Access*****/
/*****

```



```
SQL> clear scr
```

```
SQL> alter session force parallel query parallel 4;
```

```
SQL> select /*+ INMEMORY */ /*IN-MEMORY Parallel*/  
d_date,sum(l.lo_revenue) "Total Revenue"  
From   LINEORDER l, DATE_DIM d  
Where  l.lo_orderdate = d.d_datekey  
and    D_DAYNUMINMONTH = 25  
and    d.d_month = 'December'  
group by d_date  
order by d_date;
```

```
6 rows selected.
```

```
Elapsed: 00:00:02.75
```

```
Execution Plan
```

```
-----  
Plan hash value: 746328703  
  
-----  
-----  
-----
```

Id	Operation		Name	Rows	Bytes
Cost (%CPU)	Time	TQ IN-OUT PQ			
Distrib					

0	SELECT STATEMENT			7	399
755 (17)	00:00:01				

1	PX COORDINATOR				

2	PX SEND QC (ORDER)		:TQ10002	7	399
755 (17)	00:00:01	Q1,02	P->S	QC	
	(ORDER)				

3	SORT GROUP BY			7	399
755 (17)	00:00:01	Q1,02	PCWP		

4	PX RECEIVE			7	399
755 (17)	00:00:01	Q1,02	PCWP		



5	PX SEND RANGE	:TQ10001	7	399
755 (17)	00:00:01 Q1,01 P->P RA			
NGE				
6	HASH GROUP BY		7	399
755 (17)	00:00:01 Q1,01 PCWP			
* 7	HASH JOIN		7	399
754 (17)	00:00:01 Q1,01 PCWP			
* 8	TABLE ACCESS INMEMORY FULL	DATE_DIM	7	266
2 (0)	00:00:01 Q1,01 PCWP			
9	VIEW	VW_GBC_5	2407	45733
752 (17)	00:00:01 Q1,01 PCWP			
10	HASH GROUP BY		2407	28884
752 (17)	00:00:01 Q1,01 PCWP			
11	PX RECEIVE		2407	28884
752 (17)	00:00:01 Q1,01 PCWP			
12	PX SEND HASH	:TQ10000	2407	28884
752 (17)	00:00:01 Q1,00 P->P HA			
SH				
13	HASH GROUP BY		2407	28884
752 (17)	00:00:01 Q1,00 PCWP			
14	PX BLOCK ITERATOR		29M	343M
755 (17)	00:00:01 Q1,00 PCWC			
15	TABLE ACCESS INMEMORY FULL	LINEORDER	29M	343M
755 (17)	00:00:01 Q1,00 PCWP			

Predicate Information (identified by operation id):				

7 - access("ITEM_1"="D"."D_DATEKEY")				
8 - inmemory("D_DAYNUMINMONTH"=25 AND "D"."D_MONTH"='December')				



```
filter("D_DAYNUMINMONTH"=25 AND "D"."D_MONTH"='December')
```

Note

- dynamic statistics used: dynamic sampling (level=5)
- Degree of Parallelism is 4 because of session

Statistics

```
887 recursive calls
2 db block gets
1253 consistent gets
10 physical reads
804 redo size
860 bytes sent via SQL*Net to client
606 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
96 sorts (memory)
0 sorts (disk)
6 rows processed
```

Queries complejas

Three table scan

Query 1

```
SQL> conn ssb/ssb@SOE
```

```
/*****Group Aggregate with three table join*****/
/*****
/*****Parallel Disk Access*****/
```

```
SQL> clear scr
```

```
--flush the buffer_cache
```

```
SQL> alter system flush buffer_cache;
```

```
SQL> set autotrace traceonly explain statistics
```

```
SQL> alter session force parallel query parallel 4;
```

```
SQL> alter session set inmemory_query='DISABLE';
```

```
SQL> set timing on
```

```
SQL> select /*+ NO_INMEMORY */ /* DISK ACCESS */
p.p_name, sum(l.lo_revenue)
From    LINEORDER l, DATE_DIM d, PART p
Where   l.lo_orderdate = d.d_datekey
And     l.lo_partkey   = p.p_partkey
And     p.p_name       = 'misty gainsboro'
```



```

And      d.d_year = 1992
And      d.d_month = 'December'
Group by p.p_name;

```

Elapsed: 00:00:37.16

Execution Plan

Plan hash value: 2989091434

Id		Operation		Name	
Rows	Bytes	Cost (%CPU)	Time		
TQ	IN-OUT	PQ Distrib			

0	SELECT STATEMENT				
50	3400	17430 (2)	00:00:01		
1	TEMP TABLE TRANSFORMATION				
2	LOAD AS SELECT (CURSOR DURATION MEMORY)	SYS_TEMP_0FD9D6666_290D25			
3	PX COORDINATOR				
4	PX SEND QC (RANDOM)		:TQ10001		
1	24	5 (20)	00:00:01		
Q1,01	P->S	QC (RAND)			
5	HASH GROUP BY				
1	24	5 (20)	00:00:01		
Q1,01	PCWP				
6	PX RECEIVE				
1	24	4 (0)	00:00:01		
Q1,01	PCWP				
7	PX SEND HASH		:TQ10000		
1	24	4 (0)	00:00:01		
Q1,00	P->P	HASH			



8	KEY VECTOR CREATE BUFFERED	:KV0000	
1 24	4 (0) 00:00:01		
Q1,00 PCWC			
9	PX BLOCK ITERATOR		
31 620	4 (0) 00:00:01		
Q1,00 PCWC			
* 10	TABLE ACCESS FULL	DATE_DIM	
31 620	4 (0) 00:00:01		
Q1,00 PCWP			
11	LOAD AS SELECT (CURSOR DURATION MEMORY)	SYS_TEMP_0FD9D6667_290D25	
12	PX COORDINATOR		
13	PX SEND QC (RANDOM)	:TQ20001	
70 1610	397 (1) 00:00:01		
Q2,01 P->S	QC (RAND)		
14	HASH GROUP BY		
70 1610	397 (1) 00:00:01		
Q2,01 PCWP			
15	PX RECEIVE		
70 1610	396 (1) 00:00:01		
Q2,01 PCWP			
16	PX SEND HASH	:TQ20000	
70 1610	396 (1) 00:00:01		
Q2,00 P->P	HASH		
17	KEY VECTOR CREATE BUFFERED	:KV0001	
70 1610	396 (1) 00:00:01		
Q2,00 PCWC			
18	PX BLOCK ITERATOR		
71 1349	396 (1) 00:00:01		
Q2,00 PCWC			
* 19	TABLE ACCESS FULL	PART	
71 1349	396 (1) 00:00:01		
Q2,00 PCWP			
20	PX COORDINATOR		
21	PX SEND QC (RANDOM)	:TQ30003	
50 3400	17027 (2) 00:00:01		



Q3,03		P->S		QC (RAND)			
* 22				HASH JOIN BUFFERED			
50		3400		17027 (2)		00:00:01	
Q3,03		PCWP					
23				PX RECEIVE			
50		2250		17025 (2)		00:00:01	
Q3,03		PCWP					
24				PX SEND HYBRID HASH		:TQ30001	
50		2250		17025 (2)		00:00:01	
Q3,01		P->P		HYBRID HASH			
25				STATISTICS COLLECTOR			
Q3,01		PCWC					
* 26				HASH JOIN			
50		2250		17025 (2)		00:00:01	
Q3,01		PCWP					
27				TABLE ACCESS FULL		SYS_TEMP_0FD9D6666_290D25	
1		24		2 (0)		00:00:01	
Q3,01		PCWP					
28				VIEW		VW_VT_80F21617	
50		1050		17023 (2)		00:00:01	
Q3,01		PCWP					
29				HASH GROUP BY			
50		1250		17023 (2)		00:00:01	
Q3,01		PCWP					
30				PX RECEIVE			
50		1250		17023 (2)		00:00:01	
Q3,01		PCWP					
31				PX SEND HASH		:TQ30000	
50		1250		17023 (2)		00:00:01	
Q3,00		P->P		HASH			
32				VECTOR GROUP BY			
50		1250		17023 (2)		00:00:01	
Q3,00		PCWP					
33				HASH GROUP BY			
50		1250		17023 (2)		00:00:01	
Q3,00		PCWP					
34				KEY VECTOR USE		:KV0000	
70		1750		17023 (2)		00:00:01	
Q3,00		PCWC					



35		KEY VECTOR USE	:KV0001	
5417	111K	17023 (2)	00:00:01	
Q3,00	PCWC			
36		PX BLOCK ITERATOR		
29M	486M	17023 (2)	00:00:01	
Q3,00	PCWC			
* 37		TABLE ACCESS FULL	LINEORDER	
29M	486M	17023 (2)	00:00:01	
Q3,00	PCWP			
38		PX RECEIVE		
70	1610	2 (0)	00:00:01	
Q3,03	PCWP			
39		PX SEND HYBRID HASH	:TQ30002	
70	1610	2 (0)	00:00:01	
Q3,02	P->P	HYBRID HASH		
40		PX BLOCK ITERATOR		
70	1610	2 (0)	00:00:01	
Q3,02	PCWC			
41		TABLE ACCESS FULL	SYS_TEMP_0FD9D6667_290D25	
70	1610	2 (0)	00:00:01	
Q3,02	PCWP			

Predicate Information (identified by operation id):

10 - filter("D"."D_YEAR"=1992 AND "D"."D_MONTH"='December')
19 - filter("P"."P_NAME"='misty gainsboro')
22 - access("ITEM_8"=INTERNAL_FUNCTION("C0"))
26 - access("ITEM_7"=INTERNAL_FUNCTION("C0"))
37 - filter(SYS_OP_KEY_VECTOR_FILTER("L"."LO_PARTKEY",:KV0001) AND
SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000
))

Note

-
- dynamic statistics used: dynamic sampling (level=5)
 - Degree of Parallelism is 4 because of session
 - vector transformation used for this statement

Statistics




```

300 recursive calls
37 db block gets
229041 consistent gets
228675 physical reads
284504 redo size
655 bytes sent via SQL*Net to client
677 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
27 sorts (memory)
0 sorts (disk)
1 rows processed

```

```

/*****In-Memory Serial Access*****/
/*****

```

```
SQL> clear scr
```

```
SQL> alter session disable parallel query;
SQL> alter session set inmemory_query='DISABLE';
```

```
SQL> Select /*+ INMEMORY */ /*IN-MEMORY Serial*/
p.p_name, sum(l.lo_revenue)
From      LINEORDER l, DATE_DIM d, PART p
Where     l.lo_orderdate = d.d_datekey
And       l.lo_partkey   = p.p_partkey
And       p.p_name       = 'misty gainsboro'
And       d.d_year       = 1992
And       d.d_month      = 'December'
Group by p.p_name;
```

Elapsed: 00:00:00.52

Execution Plan

Plan hash value: 3307065880

Id	Operation	Name
Rows	Bytes	Cost (%CPU)
Time		
0	SELECT STATEMENT	
50	3400 2952 (22)	00:00:01
1	TEMP TABLE TRANSFORMATION	
2	LOAD AS SELECT (CURSOR DURATION MEMORY)	SYS_TEMP_0FD9D666C_290D25



3		HASH GROUP BY	
1	24	2 (50)	00:00:01
4		KEY VECTOR CREATE BUFFERED	:KV0000
1	24	1 (0)	00:00:01
* 5		TABLE ACCESS INMEMORY FULL	DATE_DIM
31	620	1 (0)	00:00:01
6		LOAD AS SELECT (CURSOR DURATION MEMORY)	SYS_TEMP_0FD9D666D_290D25
7		HASH GROUP BY	
70	1610	67 (21)	00:00:01
8		KEY VECTOR CREATE BUFFERED	:KV0001
70	1610	66 (20)	00:00:01
* 9		TABLE ACCESS INMEMORY FULL	PART
71	1349	66 (20)	00:00:01
10		HASH GROUP BY	
50	3400	2883 (22)	00:00:01
* 11		HASH JOIN	
50	3400	2882 (22)	00:00:01
12		MERGE JOIN CARTESIAN	
70	3290	4 (0)	00:00:01
13		TABLE ACCESS FULL	SYS_TEMP_0FD9D666C_290D25
1	24	2 (0)	00:00:01
14		BUFFER SORT	
70	1610	2 (0)	00:00:01
15		TABLE ACCESS FULL	SYS_TEMP_0FD9D666D_290D25
70	1610	2 (0)	00:00:01
16		VIEW	VW_VT_80F21617
50	1050	2878 (22)	00:00:01
17		VECTOR GROUP BY	
50	1250	2878 (22)	00:00:01
18		HASH GROUP BY	
50	1250	2878 (22)	00:00:01
19		KEY VECTOR USE	:KV0000
70	1750	2878 (22)	00:00:01
20		KEY VECTOR USE	:KV0001
5417	111K	2878 (22)	00:00:01
* 21		TABLE ACCESS INMEMORY FULL	LINEORDER
29M	486M	2878 (22)	00:00:01

Predicate Information (identified by operation id):

```

5 - inmemory("D"."D_YEAR"=1992 AND "D"."D_MONTH"='December')
    filter("D"."D_YEAR"=1992 AND "D"."D_MONTH"='December')
9 - inmemory("P"."P_NAME"='misty gainsboro')
    filter("P"."P_NAME"='misty gainsboro')
11 - access("ITEM_7"=INTERNAL_FUNCTION("C0") AND
"ITEM_8"=INTERNAL_FUNCTION("C0"))
21 - inmemory(SYS_OP_KEY_VECTOR_FILTER("L"."LO_PARTKEY",:KV0001) AND
SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000))
    filter(SYS_OP_KEY_VECTOR_FILTER("L"."LO_PARTKEY",:KV0001) AND
SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000))

```



Note

- vector transformation used for this statement

Statistics

```
21 recursive calls
0 db block gets
38 consistent gets
0 physical reads
0 redo size
655 bytes sent via SQL*Net to client
677 bytes received via SQL*Net from client
2 SQL*Net roundtrips to/from client
2 sorts (memory)
0 sorts (disk)
1 rows processed
```

```
/******In-Memory Parallel Access*****/
/******
```

```
SQL> clear scr
```

```
SQL> alter session force parallel query parallel 4;
```

```
SQL> select /*+ INMEMORY */ /*IN-MEMORY Parallel*/
p.p_name, sum(l.lo_revenue)
From      LINEORDER l, DATE_DIM d, PART p
Where     l.lo_orderdate = d.d_datekey
And       l.lo_partkey   = p.p_partkey
And       p.p_name       = 'misty gainsboro'
And       d.d_year       = 1992
And       d.d_month      = 'December'
Group by p.p_name;
```

```
Elapsed: 00:00:00.44
```

Execution Plan

```
-----
Plan hash value: 2989091434
-----
-----
-----
```

Id	Operation	Name
Rows	Bytes Cost (%CPU) Time	



TQ	IN-OUT	PQ Distrib	

0	SELECT STATEMENT		
50	3400	827 (22)	00:00:01
1	TEMP TABLE TRANSFORMATION		
2	LOAD AS SELECT (CURSOR DURATION MEMORY)		
			SYS_TEMP_0FD9D6672_290D25
3	PX COORDINATOR		
4	PX SEND QC (RANDOM)		
1	24	3 (34)	00:00:01
Q1,01	P->S	QC (RAND)	:TQ10001
5	HASH GROUP BY		
1	24	3 (34)	00:00:01
Q1,01	PCWP		
6	PX RECEIVE		
1	24	2 (0)	00:00:01
Q1,01	PCWP		
7	PX SEND HASH		
1	24	2 (0)	00:00:01
Q1,00	P->P	HASH	:TQ10000
8	KEY VECTOR CREATE BUFFERED		
1	24	2 (0)	00:00:01
Q1,00	PCWC		:KV0000
9	PX BLOCK ITERATOR		
31	620	2 (0)	00:00:01
Q1,00	PCWC		
* 10	TABLE ACCESS INMEMORY FULL		
31	620	2 (0)	00:00:01
Q1,00	PCWP		DATE_DIM
11	LOAD AS SELECT (CURSOR DURATION MEMORY)		
			SYS_TEMP_0FD9D6673_290D25



12	PX COORDINATOR	
13	PX SEND QC (RANDOM)	:TQ20001
70	1610	19 (22) 00:00:01
Q2,01	P->S	QC (RAND)
14	HASH GROUP BY	
70	1610	19 (22) 00:00:01
Q2,01	PCWP	
15	PX RECEIVE	
70	1610	18 (17) 00:00:01
Q2,01	PCWP	
16	PX SEND HASH	:TQ20000
70	1610	18 (17) 00:00:01
Q2,00	P->P	HASH
17	KEY VECTOR CREATE BUFFERED	:KV0001
70	1610	18 (17) 00:00:01
Q2,00	PCWC	
18	PX BLOCK ITERATOR	
71	1349	18 (17) 00:00:01
Q2,00	PCWC	
* 19	TABLE ACCESS INMEMORY FULL	PART
71	1349	18 (17) 00:00:01
Q2,00	PCWP	
20	PX COORDINATOR	
21	PX SEND QC (RANDOM)	:TQ30003
50	3400	803 (22) 00:00:01
Q3,03	P->S	QC (RAND)
* 22	HASH JOIN BUFFERED	
50	3400	803 (22) 00:00:01
Q3,03	PCWP	
23	PX RECEIVE	
50	2250	801 (22) 00:00:01
Q3,03	PCWP	
24	PX SEND HYBRID HASH	:TQ30001
50	2250	801 (22) 00:00:01
Q3,01	P->P	HYBRID HASH
25	STATISTICS COLLECTOR	



	Q3,01		PCWC			
	* 26				HASH JOIN	
	50		2250		801 (22)	00:00:01
	Q3,01		PCWP			
	27				TABLE ACCESS FULL	SYS_TEMP_0FD9D6672_290D25
	1		24		2 (0)	00:00:01
	Q3,01		PCWP			
	28				VIEW	VW_VT_80F21617
	50		1050		799 (22)	00:00:01
	Q3,01		PCWP			
	29				HASH GROUP BY	
	50		1250		799 (22)	00:00:01
	Q3,01		PCWP			
	30				PX RECEIVE	
	50		1250		799 (22)	00:00:01
	Q3,01		PCWP			
	31				PX SEND HASH	:TQ30000
	50		1250		799 (22)	00:00:01
	Q3,00		P->P		HASH	
	32				VECTOR GROUP BY	
	50		1250		799 (22)	00:00:01
	Q3,00		PCWP			
	33				HASH GROUP BY	
	50		1250		799 (22)	00:00:01
	Q3,00		PCWP			
	34				KEY VECTOR USE	:KV0000
	70		1750		799 (22)	00:00:01
	Q3,00		PCWC			
	35				KEY VECTOR USE	:KV0001
	5417		111K		799 (22)	00:00:01
	Q3,00		PCWC			
	36				PX BLOCK ITERATOR	
	29M		486M		799 (22)	00:00:01
	Q3,00		PCWC			
	* 37				TABLE ACCESS INMEMORY FULL	LINEORDER
	29M		486M		799 (22)	00:00:01
	Q3,00		PCWP			
	38				PX RECEIVE	
	70		1610		2 (0)	00:00:01
	Q3,03		PCWP			



```

| 39 |          PX SEND HYBRID HASH          | :TQ30002
| 70 | 1610 |      2  (0) | 00:00:01
| Q3,02 | P->P | HYBRID HASH|

| 40 |          PX BLOCK ITERATOR          |
| 70 | 1610 |      2  (0) | 00:00:01
| Q3,02 | PCWC |          |

| 41 |          TABLE ACCESS FULL          | SYS_TEMP_0FD9D6673_290D25
| 70 | 1610 |      2  (0) | 00:00:01
| Q3,02 | PCWP |          |

```

```

-----
-----
-----

```

Predicate Information (identified by operation id):

```

-----
10 - inmemory("D"."D_YEAR"=1992 AND "D"."D_MONTH"='December')
      filter("D"."D_YEAR"=1992 AND "D"."D_MONTH"='December')
19 - inmemory("P"."P_NAME"='misty gainsboro')
      filter("P"."P_NAME"='misty gainsboro')
22 - access("ITEM_8"=INTERNAL_FUNCTION("C0"))
26 - access("ITEM_7"=INTERNAL_FUNCTION("C0"))
37 - inmemory(SYS_OP_KEY_VECTOR_FILTER("L"."LO_PARTKEY",:KV0001) AND
SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV00
00))

      filter(SYS_OP_KEY_VECTOR_FILTER("L"."LO_PARTKEY",:KV0001) AND
SYS_OP_KEY_VECTOR_FILTER("L"."LO_ORDERDATE",:KV0000
))

```

Note

```

-----
- dynamic statistics used: dynamic sampling (level=5)
- Degree of Parallelism is 4 because of session
- vector transformation used for this statement

```

Statistics

```

-----
108 recursive calls
  0 db block gets
162 consistent gets
  0 physical reads
  0 redo size
655 bytes sent via SQL*Net to client
679 bytes received via SQL*Net from client
  2 SQL*Net roundtrips to/from client
  8 sorts (memory)
  0 sorts (disk)

```



1 rows processed

ACO - Oracle Advanced Compression (45 min)

Oracle Advanced Compression (ACO) permite aumentar el rendimiento al mismo tiempo que se reduce el coste del almacenamiento. Permite reducir significativamente el espacio total de almacenamiento de la base de datos al permitir la compresión para todo tipo de datos relacionales (tabla), no estructurados (archivo), índice, red, Data Guard, backup con RMAN.

Mediante *Advanced Row Compression* se puede reducir el consumo de almacenamiento en un factor de 2x a 4x. De igual forma, los accesos por parte del optimizador, se incrementan en 2,5x en procesos de table-scan comparados con los datos no comprimidos.

Los datos se leen comprimidos (datos e índices) directamente, en memoria, sin descomprimir los bloques.

Advanced Row Compression Implementation

Se utilizará el usuario SSB donde se crearán dos tablas

- TEST_TAB_NO (sin compresión)
- TEST_TAB_COMP (con compression)

Ejecutaremos inserciones en ambas tablas para comprobar el factor de compresión obtenido mediante el uso del API (DBMS_COMPRESSION) y comprobaremos los bloques de ambas tablas.

Para ello utilizaremos la PDB (SOE) y nos conectaremos con usuario SSB.



Crear dos tablas: TEST_TAB_NO y TEST_TAB_COMP y se procederá a insertar registros sobre la tabla TEST_TAB_NO que no está comprimida.

```
[oracle@single19c ~]$ sqlplus ssb/ssb@soe

CREATE TABLE test_tab_no (
  id          NUMBER(10)    NOT NULL,
  description  VARCHAR2(100) NOT NULL,
  created_date DATE         NOT NULL,
  created_by   VARCHAR2(50)  NOT NULL,
  updated_date DATE,
  updated_by   VARCHAR2(50)
);

Table created.

SQL> SET SERVEROUTPUT ON

SQL> DECLARE
  v_date      test_tab_no.created_date%TYPE := SYSDATE;
  v_user      test_tab_no.created_by%TYPE   := USER;
  l_start_time NUMBER;
  l_start_cpu  NUMBER;
BEGIN

  l_start_time := DBMS_UTILITY.get_time;
  l_start_cpu  := DBMS_UTILITY.get_cpu_time;

  INSERT /*+ APPEND */ INTO test_tab_no (id, description, created_date,
created_by)
  SELECT level,
         'A very repetitive, and therefore very compressible column value',
         v_date,
         v_user
  FROM   dual
  CONNECT BY level <= 1000000;
  COMMIT;

  DBMS_OUTPUT.put_line('CPU Time (hsecs)      : ' || (DBMS_UTILITY.get_cpu_time -
l_start_cpu));
  DBMS_OUTPUT.put_line('Elapsed Time (hsecs): ' || (DBMS_UTILITY.get_time -
l_start_time));
END;
/

CPU Time (hsecs)      : 535
Elapsed Time (hsecs): 702

PL/SQL procedure successfully completed.
```



Utilizaremos el API DBMS_COMPRESSION con el tipo de compresión avanzada (DBMS_COMPRESSION.comp_advanced) y veremos el factor de compresión que se podría lograr en esta tabla no comprimida.

```
SET SERVEROUTPUT ON

DECLARE
    l_blkcnt_cmp      PLS_INTEGER;
    l_blkcnt_uncomp   PLS_INTEGER;
    l_row_cmp         PLS_INTEGER;
    l_row_uncomp      PLS_INTEGER;
    l_cmp_ratio       NUMBER;
    l_comptype_str    VARCHAR2(32767);
BEGIN
    DBMS_COMPRESSION.get_compression_ratio (
        scratchtbsname => 'SYSAUX',
        ownname         => 'SSB',
        objname         => 'TEST_TAB_NO',
        subobjname      => NULL,
        comptype        => DBMS_COMPRESSION.comp_advanced,
        blkcnt_cmp      => l_blkcnt_cmp,
        blkcnt_uncomp   => l_blkcnt_uncomp,
        row_cmp         => l_row_cmp,
        row_uncomp      => l_row_uncomp,
        cmp_ratio       => l_cmp_ratio,
        comptype_str    => l_comptype_str,
        subset_numrows  => DBMS_COMPRESSION.comp_ratio_allrows,
        objtype         => DBMS_COMPRESSION.objtype_table
    );

    DBMS_OUTPUT.put_line('Number of blocks used (compressed)      : ' ||
l_blkcnt_cmp);
    DBMS_OUTPUT.put_line('Number of blocks used (uncompressed)   : ' ||
l_blkcnt_uncomp);
```



```

    DBMS_OUTPUT.put_line('Number of rows in a block (compressed)   : ' ||
l_row_cmp);
    DBMS_OUTPUT.put_line('Number of rows in a block (uncompressed) : ' ||
l_row_uncmp);
    DBMS_OUTPUT.put_line('Compression ratio                        : ' ||
l_cmp_ratio);
    DBMS_OUTPUT.put_line('Compression type                        : ' ||
l_comptype_str);
END;
/
Number of blocks used (compressed)      : 216
Number of blocks used (uncompressed)    : 1664
Number of rows in a block (compressed)  : 647
Number of rows in a block (uncompressed): 84
Compression ratio                       : 7.7
Compression type                        : "Compress Advanced"

PL/SQL procedure successfully completed.

```

Crear la tabla TEST_TAB_COMP para insertar registros sobre la tabla comprimida.
Utilizaremos el atributo COMPRESS FOR ALL OPERATIONS en la creación de la tabla.

```

SQL> CREATE TABLE test_tab_comp (
    id          NUMBER(10)    NOT NULL,
    description  VARCHAR2(100) NOT NULL,
    created_date DATE          NOT NULL,
    created_by   VARCHAR2(50)  NOT NULL,
    updated_date DATE,
    updated_by   VARCHAR2(50)
)
COMPRESS FOR ALL OPERATIONS;
Table created.

SQL> SET SERVEROUTPUT ON
SQL> DECLARE

```



```

v_date      test_tab_comp.created_date%TYPE := SYSDATE;
v_user      test_tab_comp.created_by%TYPE   := USER;
l_start_time NUMBER;
l_start_cpu  NUMBER;
BEGIN

    l_start_time := DBMS_UTILITY.get_time;
    l_start_cpu  := DBMS_UTILITY.get_cpu_time;

    INSERT /*+ APPEND */ INTO test_tab_comp (id, description, created_date,
created_by)
    SELECT level,
           'A very repetitive, and therefore very compressible column value',
           v_date,
           v_user
    FROM    dual
    CONNECT BY level <= 1000000;
    COMMIT;

    DBMS_OUTPUT.put_line('CPU Time (hsecs)      : ' || (DBMS_UTILITY.get_cpu_time -
l_start_cpu));
    DBMS_OUTPUT.put_line('Elapsed Time (hsecs): ' || (DBMS_UTILITY.get_time -
l_start_time));
END;
/
CPU Time (hsecs)      : 349
Elapsed Time (hsecs): 370

PL/SQL procedure successfully completed.

```

Comprobar la ocupación de bloques de ambas tablas para ver la mejora con la compresión.

```
col table_name format a20
```



```
SQL> SELECT table_name,compression, compress_for
FROM   user_tables
WHERE  table_name like 'TEST_TAB%';
```

TABLE_NAME	COMPRESS	COMPRESS_FOR
TEST_TAB_COMP	ENABLED	ADVANCED
TEST_TAB_NO	DISABLED	

```
SQL> SELECT table_name,
           compression,
           num_rows,
           blocks,
           empty_blocks
FROM   user_tables
WHERE  table_name like 'TEST_TAB%'
ORDER BY 1;
```

TABLE_NAME	COMPRESS	NUM_ROWS	BLOCKS	EMPTY_BLOCKS
TEST_TAB_COMP	ENABLED	1000000	1578	0
TEST_TAB_NO	DISABLED	1000000	12055	0

Advanced Row Compression Tabla/Partición

De igual forma, se podría comprimir particiones de tablas particionadas para implementar si se quisiera un entorno ILM (*Information Lifecycle Management*) donde podríamos definir políticas sobre las particiones y así poder almacenar información que va a ser accedida frecuentemente o cada cierto tiempo.

Veamos un ejemplo de compresión a nivel de partición. Se creará una tabla con dos particiones y procederemos ejecutar varios procesos DML (Data Manipulation Language) para insertar información sobre las particiones creadas.



La tabla será creada para comprimir sólo datos sobre una partición definida para compresión.

```
SQL> CREATE TABLE test_tab_particion (  
  id          NUMBER(10)    NOT NULL,  
  description VARCHAR2(100) NOT NULL,  
  created_date DATE          NOT NULL,  
  created_by   VARCHAR2(50)  NOT NULL,  
  updated_date DATE,  
  updated_by   VARCHAR2(50)  
)  
NOCOMPRESS  
PARTITION BY RANGE (created_date) (  
  PARTITION test_tab_q1 VALUES LESS THAN (TO_DATE('01/04/2003', 'DD/MM/YYYY'))  
  COMPRESS FOR ALL OPERATIONS,  
  PARTITION test_tab_q2 VALUES LESS THAN (MAXVALUE)  
);
```

Table created

```
SQL> set linesize 1000;
```

```
SQL> COLUMN partition_name FORMAT A30
```

```
SQL> SELECT partition_name, compression, compress_for  
FROM   user_tab_partitions  
WHERE  table_name = 'TEST_TAB_PARTICION'  
ORDER BY 1;
```

PARTITION_NAME	COMPRESS	COMPRESS_FOR
TEST_TAB_Q1	ENABLED	ADVANCED
TEST_TAB_Q2	DISABLED	



Ejecutaremos dos procesos DML para insertar datos de varias fechas repartidos en las dos particiones creadas, una con compresión y otra sin ella, y simular una implementación de tablas tipo ILM. En primer lugar sólo se inserta en la partición sin compresión.

```
SQL> SET SERVEROUTPUT ON
SQL> DECLARE
    v_date          test_tab_particion.created_date%TYPE := SYSDATE;
    v_user          test_tab_particion.created_by%TYPE   := USER;
    l_start_time    NUMBER;
    l_start_cpu     NUMBER;
BEGIN

    l_start_time := DBMS_UTILITY.get_time;
    l_start_cpu  := DBMS_UTILITY.get_cpu_time;

    INSERT /*+ APPEND */ INTO test_tab_particion (id, description, created_date,
created_by)
    SELECT level,
           'A very repetitive, and therefore very compressible column value',
           v_date,
           v_user
    FROM    dual
    CONNECT BY level <= 1000000;
    COMMIT;

    DBMS_OUTPUT.put_line('CPU Time (hsecs)      : ' || (DBMS_UTILITY.get_cpu_time -
l_start_cpu));
    DBMS_OUTPUT.put_line('Elapsed Time (hsecs): ' || (DBMS_UTILITY.get_time -
l_start_time));
END;
/
CPU Time (hsecs)      : 359
Elapsed Time (hsecs): 382
```



```
PL/SQL procedure successfully completed.
```

Analizar la tabla (Partición) mediante DBMS_STATS.gather_table_stats. Posteriormente veremos la ocupación de los bloques y datos insertados sobre la tabla particionada.

```
SQL> EXEC DBMS_STATS.gather_table_stats(USER, 'TEST_TAB_PARTICION');
PL/SQL procedure successfully completed.
```

```
SQL> SELECT table_name,
           partition_name,
           compression,
           num_rows,
           blocks,
           empty_blocks
FROM   user_tab_partitions
WHERE  table_name = 'TEST_TAB_PARTICION'
ORDER BY 1;
```

TABLE_NAME	PARTITION_NAME	COMPRESS	NUM_ROWS
BLOCKS	EMPTY_BLOCKS		

TEST_TAB_PARTICION	TEST_TAB_Q1	ENABLED	0
0	0		
TEST_TAB_PARTICION	TEST_TAB_Q2	DISABLED	1000000
12038	0		

Repetimos el proceso de inserción en la tabla particionada por diferentes fechas para crear registros en la partición definida con el atributo COMPRESS FOR ALL OPERATIONS.

```
SQL> SET SERVEROUTPUT ON
```

```
SQL> DECLARE
  v_date      test_tab_particion.created_date%TYPE := TO_DATE('31/03/2003',
'DD/MM/YYYY');
  v_user      test_tab_particion.created_by%TYPE   := USER;
  l_start_time NUMBER;
```




```

    l_start_cpu    NUMBER;
BEGIN

    l_start_time := DBMS_UTILITY.get_time;
    l_start_cpu  := DBMS_UTILITY.get_cpu_time;

    INSERT /*+ APPEND */ INTO test_tab_particion (id, description, created_date,
created_by)
    SELECT level,
           'A very repetitive, and therefore very compressible column value',
           v_date,
           v_user
    FROM    dual
    CONNECT BY level <= 1000000;
    COMMIT;

    DBMS_OUTPUT.put_line('CPU Time (hsecs)      : ' || (DBMS_UTILITY.get_cpu_time -
l_start_cpu));
    DBMS_OUTPUT.put_line('Elapsed Time (hsecs): ' || (DBMS_UTILITY.get_time -
l_start_time));
END;
/

CPU Time (hsecs)      : 329
Elapsed Time (hsecs): 391

PL/SQL procedure successfully completed.

```

Analizar la tabla (Partición) mediante DBMS_STATS.gather_table_stats. Posteriormente veremos la ocupación de los bloques y datos insertados sobre la tabla particionada. Como vemos, en la partición comprimida hay menos bloques con respecto a la partición sin comprimir

```

SQL> EXEC DBMS_STATS.gather_table_stats(USER, 'TEST_TAB_PARTICION');

PL/SQL procedure successfully completed.

SQL> COLUMN table_name FORMAT A20
SQL> COLUMN partition_name FORMAT A20

SQL> SELECT table_name,
           partition_name,
           compression,
           num_rows,
           blocks,
           empty_blocks
    FROM    user_tab_partitions
   WHERE   table_name = 'TEST_TAB_PARTICION'
  ORDER BY 1;

```



TABLE_NAME EMPTY_BLOCKS	PARTITION_NAME	COMPRESS	NUM_ROWS	BLOCKS
-----	-----	-----	-----	-----
TEST_TAB_PARTICION 0	TEST_TAB_Q1	ENABLED	1000000	1578
TEST_TAB_PARTICION 0	TEST_TAB_Q2	DISABLED	1000000	12038

Utilización de API PL/SQL DBMS_COMPRESSION

A continuación, veremos los distintos usos que podemos lograr con el API PL/SQL DBMS_COMPRESSION utilizando distintos atributos de tipo de compresión. Para una lista completa de todos los valores soportados consultar el manual de esta API.

Para ello, crearemos una tabla particionada e insertaremos y chequearemos los distintos factores de compresión.

Posteriormente, se creará un índice local sobre la tabla TAB_DBMS_COMPRESS.

El procedimiento GET_COMPRESSION_RATIO estima el impacto de diferentes niveles de compresión en una tabla o partición especificada.

```
SQL> CREATE TABLE tab_dbms_compress (
  id          NUMBER,
  code        VARCHAR2(20),
  description  VARCHAR2(50),
  clob_description CLOB,
  created_date DATE,
  CONSTRAINT tab_dbms_compress_pk PRIMARY KEY (id)
)
PARTITION BY RANGE (created_date)
(PARTITION tab_dbms_compress_part_2015 VALUES LESS THAN (TO_DATE('01/01/2016',
'DD/MM/YYYY')) TABLESPACE sysaux,
 PARTITION tab_dbms_compress_part_2016 VALUES LESS THAN (TO_DATE('01/01/2017',
'DD/MM/YYYY')) TABLESPACE sysaux);

Table created.

--- Creamos un indice sobre la tabla tab_dbms_compress ---

SQL> CREATE INDEX tab_dbms_compress_code_idx ON tab_dbms_compress(code) LOCAL;

Index created

SQL> INSERT INTO tab_dbms_compress
```



```

SELECT level,
       CASE
         WHEN MOD(level,2)=0 THEN 'CODE1'
         ELSE 'CODE2'
       END,
       CASE
         WHEN MOD(level,2)=0 THEN 'Description for CODE1'
         ELSE 'Description for CODE2'
       END,
       CASE
         WHEN MOD(level,2)=0 THEN 'CLOB description for CODE1'
         ELSE 'CLOB description for CODE2'
       END,
       CASE
         WHEN MOD(level,2)=0 THEN TO_DATE('01/07/2015','DD/MM/YYYY')
         ELSE TO_DATE('01/07/2016','DD/MM/YYYY')
       END
FROM   dual
CONNECT BY level <= 100000;
COMMIT;

100000 rows created.

Commit complete.

SQL> EXEC DBMS_STATS.gather_table_stats(USER, 'tab_dbms_compress');

PL/SQL procedure successfully completed.

```

El primer ejemplo muestra el efecto de la compresión OLTP en una tabla específica, utilizando todas las filas de la tabla como tamaño de muestra.

```

SQL> SET SERVEROUTPUT ON

SQL> DECLARE
  l_blkcnt_cmp      PLS_INTEGER;
  l_blkcnt_uncmp    PLS_INTEGER;
  l_row_cmp         PLS_INTEGER;
  l_row_uncmp       PLS_INTEGER;
  l_cmp_ratio       NUMBER;
  l_comptype_str    VARCHAR2(32767);
BEGIN
  DBMS_COMPRESSION.get_compression_ratio (
    scratchtbsname => 'SYSAUX',
    ownname        => 'SSB',
    objname        => 'TAB_DBMS_COMPRESS',
    subobjname     => NULL,
    comptype       => DBMS_COMPRESSION.comp_advanced,
    blkcnt_cmp     => l_blkcnt_cmp,
    blkcnt_uncmp   => l_blkcnt_uncmp,
    row_cmp        => l_row_cmp,
    row_uncmp      => l_row_uncmp,
    cmp_ratio      => l_cmp_ratio,

```



```

    comptype_str    => l_comptype_str,
    subset_numrows  => DBMS_COMPRESSION.comp_ratio_allrows,
    objtype         => DBMS_COMPRESSION.objtype_table
);

DBMS_OUTPUT.put_line('Number of blocks used (compressed)      : ' ||
l_blkcnt_cmp);
DBMS_OUTPUT.put_line('Number of blocks used (uncompressed)    : ' ||
l_blkcnt_uncmp);
DBMS_OUTPUT.put_line('Number of rows in a block (compressed)  : ' ||
l_row_cmp);
DBMS_OUTPUT.put_line('Number of rows in a block (uncompressed): ' ||
l_row_uncmp);
DBMS_OUTPUT.put_line('Compression ratio                      : ' ||
l_cmp_ratio);
DBMS_OUTPUT.put_line('Compression type                        : ' ||
l_comptype_str);
END;
/

Number of blocks used (compressed)      : 1313
Number of blocks used (uncompressed)    : 1737
Number of rows in a block (compressed)  : 74
Number of rows in a block (uncompressed): 55
Compression ratio                      : 1.3
Compression type                        : "Compress Advanced"

PL/SQL procedure successfully completed.

```

El segundo ejemplo muestra el efecto de la compresión OLTP en un índice particionado, utilizando todas las filas de la tabla como tamaño de muestra.

```

SQL> SET SERVEROUTPUT ON

SQL> DECLARE
    l_blkcnt_cmp    PLS_INTEGER;
    l_blkcnt_uncmp  PLS_INTEGER;
    l_row_cmp       PLS_INTEGER;
    l_row_uncmp     PLS_INTEGER;
    l_cmp_ratio     NUMBER;
    l_comptype_str  VARCHAR2(32767);
BEGIN
    DBMS_COMPRESSION.get_compression_ratio (
        scratchtbsname => 'SYSAUX',
        ownname        => 'SSB',
        objname         => 'TAB_DBMS_COMPRESS_CODE_IDX',
        subobjname      => 'TAB_DBMS_COMPRESS_PART_2015',
        comptype        => DBMS_COMPRESSION.comp_index_advanced_low,
        blkcnt_cmp      => l_blkcnt_cmp,
        blkcnt_uncmp    => l_blkcnt_uncmp,
        row_cmp         => l_row_cmp,
        row_uncmp       => l_row_uncmp,

```



```

        cmp_ratio      => l_cmp_ratio,
        comptype_str   => l_comptype_str,
        subset_numrows => DBMS_COMPRESSION.comp_ratio_minrows,
        objtype        => DBMS_COMPRESSION.objtype_index
    );

    DBMS_OUTPUT.put_line('Number of blocks used (compressed)      : ' ||
l_blkcnt_cmp);
    DBMS_OUTPUT.put_line('Number of blocks used (uncompressed)   : ' ||
l_blkcnt_uncmp);
    DBMS_OUTPUT.put_line('Number of rows in a block (compressed) : ' ||
l_row_cmp);
    DBMS_OUTPUT.put_line('Number of rows in a block (uncompressed) : ' ||
l_row_uncmp);
    DBMS_OUTPUT.put_line('Compression ratio                      : ' ||
l_cmp_ratio);
    DBMS_OUTPUT.put_line('Compression type                      : ' ||
l_comptype_str);
END;
/

Number of blocks used (compressed)      : 78
Number of blocks used (uncompressed)   : 120
Number of rows in a block (compressed) : 641
Number of rows in a block (uncompressed) : 417
Compression ratio                      : 1.5
Compression type                      : "Compress Advanced Low"

PL/SQL procedure successfully completed.

```

El tercer ejemplo muestra el efecto de la compresión OLTP en un campo LOB, utilizando todas las filas de la tabla como tamaño de muestra.

```

SQL> SET SERVEROUTPUT ON

SQL> DECLARE
    l_blkcnt_cmp      PLS_INTEGER;
    l_blkcnt_uncmp    PLS_INTEGER;
    l_lobcnt          PLS_INTEGER;
    l_cmp_ratio       NUMBER;
    l_comptype_str    VARCHAR2(32767);
BEGIN
    DBMS_COMPRESSION.get_compression_ratio (
        scratchtbsname => 'SYSAUX',
        tabowner       => 'SSB',
        tabname        => 'TAB_DBMS_COMPRESS',
        lobname        => 'CLOB_DESCRIPTION',
        partname       => NULL,
        comptype       => DBMS_COMPRESSION.comp_lob_high,
        blkcnt_cmp     => l_blkcnt_cmp,
        blkcnt_uncmp   => l_blkcnt_uncmp,
        lobcnt         => l_lobcnt,

```



```

        cmp_ratio      => l_cmp_ratio,
        comptype_str   => l_comptype_str,
        subset_numrows => DBMS_COMPRESSION.comp_ratio_lob_maxrows
    );

    DBMS_OUTPUT.put_line('Number of blocks used (compressed)      : ' ||
l_blkcnt_cmp);
    DBMS_OUTPUT.put_line('Number of blocks used (uncompressed)   : ' ||
l_blkcnt_uncmp);
    DBMS_OUTPUT.put_line('Number of rows in a block (compressed) : ' ||
l_lobcnt);
    DBMS_OUTPUT.put_line('Number of lobes sampled                : ' ||
l_cmp_ratio);
    DBMS_OUTPUT.put_line('Compression type                      : ' ||
l_comptype_str);
END;
/

Number of blocks used (compressed)      : 68
Number of blocks used (uncompressed)   : 61
Number of rows in a block (compressed) : 4999
Number of lobes sampled                : .8
Compression type                      : "Compress High"

PL/SQL procedure successfully completed.

```

La función GET_COMPRESSION_TYPE muestra el nivel de compresión para una fila especificada en una tabla. En este caso concreto, al tratarse de una tabla que no está comprimida la salida será 'COMP_NOCOMPRESS'.

```

SQL> SELECT rowid,
        CASE DBMS_COMPRESSION.get_compression_type ('SSB', 'TAB_DBMS_COMPRESS',
rowid, 'TAB_DBMS_COMPRESS_PART_2015')
        WHEN 1      THEN 'COMP_NOCOMPRESS'
        WHEN 2      THEN 'COMP_ADVANCED'
        WHEN 4      THEN 'COMP_QUERY_HIGH'
        WHEN 8      THEN 'COMP_QUERY_LOW'
        WHEN 16     THEN 'COMP_ARCHIVE_HIGH'
        WHEN 32     THEN 'COMP_ARCHIVE_LOW'
        WHEN 64     THEN 'COMP_BLOCK'
        WHEN 128    THEN 'COMP_LOB_HIGH'
        WHEN 256    THEN 'COMP_LOB_MEDIUM'
        WHEN 512    THEN 'COMP_LOB_LOW'
        WHEN 1024   THEN 'COMP_INDEX_ADVANCED_HIGH'
        WHEN 2048   THEN 'COMP_INDEX_ADVANCED_LOW'
        WHEN 1000   THEN 'COMP_RATIO_LOB_MINROWS'
        WHEN 4096   THEN 'COMP_BASIC'
        WHEN 5000   THEN 'COMP_RATIO_LOB_MAXROWS'
        WHEN 8192   THEN 'COMP_INMEMORY_NOCOMPRESS'
        WHEN 16384  THEN 'COMP_INMEMORY_DML'
        WHEN 32768  THEN 'COMP_INMEMORY_QUERY_LOW'
        WHEN 65536  THEN 'COMP_INMEMORY_QUERY_HIGH'

```



```

        WHEN 32768 THEN 'COMP_INMEMORY_CAPACITY_LOW'
        WHEN 65536 THEN 'COMP_INMEMORY_CAPACITY_HIGH'
    END AS compression_type
FROM    ssb.TAB_DBMS_COMPRESS PARTITION (TAB_DBMS_COMPRESS_part_2015)
WHERE   rownum <= 5;
/

ROWID                COMPRESSION_TYPE
-----
AAASYdAAEAAAMwDAAA  COMP_NOCOMPRESS
AAASYdAAEAAAMwDAAB  COMP_NOCOMPRESS
AAASYdAAEAAAMwDAAC  COMP_NOCOMPRESS
AAASYdAAEAAAMwDAAD  COMP_NOCOMPRESS
AAASYdAAEAAAMwDAAE  COMP_NOCOMPRESS

```

Proceso de compresión de índice

De igual forma que se puede comprimir una tabla/partición, también es importante poder comprimir los índices aplicados sobre los bloques de las tablas.

En el siguiente ejemplo utilizaremos distintos factores de compresión para los índices bien sean únicos/no únicos creados en la tabla TEST.

```

SQL> conn ssb/ssb@soe

SQL> CREATE TABLE test (
ENAME VARCHAR2(75),
EADD1 VARCHAR2(75),
EADD2 VARCHAR2(75),
EADD3 VARCHAR2(75),
EADD4 VARCHAR2(75),
CITY VARCHAR2(75)
);

Table created.

```

Ahora insertaremos datos mediante insert con hint /*+ APPEND */ sobre la tabla TEST

```

SQL> INSERT /*+ APPEND */ INTO test
SELECT RPAD('X',75, 'X'),
RPAD('X',75, 'X'),
RPAD('X',75, 'X'),
RPAD('X',75, 'X'),
RPAD('X',75, 'X'),
RPAD(TO_CHAR(level),75, 'X')
FROM dual
CONNECT BY level <= 10000;
COMMIT;

```



```
10000 rows created.
```

```
Commit complete.
```

```
col owner format a10
col segment_name format a15
col segment_type format a15
```

```
SQL> SELECT table_name, blocks
FROM dba_tables
WHERE table_name IN ('TEST')
/
```

TABLE_NAME	BLOCKS
TEST	687

Creamos un primer índice no único sin compresión ('TEST_IDX'), verificamos su tamaño y lo borramos. Estas operaciones las vamos a realizar como sys para tener acceso a determinadas vistas del diccionario de datos que requieren privilegios de sysdba.

```
SQL> conn sys/we1c0m3_we1c0m3_@SOE as sysdba
```

```
Connected.
```

```
SQL> CREATE INDEX ssb.test_idx ON ssb.test(ENAME, EADD1, EADD2, EADD3, EADD4,
CITY);
```

```
Index created
```

```
SQL> select owner,segment_name,segment_type,bytes from dba_segments where
owner='SSB' and segment_name in ('TEST','TEST_IDX');
```

OWNER	SEGMENT_NAME	SEGMENT_TYPE	BYTES
SSB	TEST	TABLE	6291456
SSB	TEST_IDX	INDEX	6291456

```
SQL> drop index ssb.TEST_IDX;
```

Creamos un segundo índice no único CON compresión ('TEST_IDX') y verificamos su tamaño. Para ello, se usará el atributo COMPRESS en la creación del índice. Después lo borramos.

```
SQL> CREATE INDEX ssb.test_idx ON ssb.test(ENAME, EADD1, EADD2, EADD3, EADD4,
CITY) COMPRESS 5;
```

```
SQL> select owner,segment_name,segment_type,bytes from dba_segments where
owner='SSB' and segment_name in ('TEST','TEST_IDX');
```



OWNER	SEGMENT_NAME	SEGMENT_TYPE	BYTES
SSB	TEST	TABLE	6291456
SSB	TEST_IDX	INDEX	2097152

```
drop index ssb.test_idx;
```

Por último creamos un índice no único con compresión COMPRESS ADVANCED LOW y COMPRESS ADVANCED HIGH verificando el tamaño de este índice 'TEST_IDX'. Como podemos comprobar la compresión tipo COMPRESS ADVANCED HIGH comprime los índices de una forma más óptima. El tipo de compresión HIGH está disponible a partir de la versión 12.2 de Oracle Database.

```
SQL> CREATE UNIQUE INDEX ssb.test_idx ON ssb.test(ENAME, EADD1, EADD2, EADD3, EADD4, CITY) COMPRESS ADVANCED LOW;
```

```
SQL> select owner,segment_name,segment_type,bytes from dba_segments where owner='SSB' and segment_name in ('TEST','TEST_IDX');
```

OWNER	SEGMENT_NAME	SEGMENT_TYPE	BYTES
SSB	TEST	TABLE	6291456
SSB	TEST_IDX	INDEX	2097152

```
drop index ssb.test_idx;
```

```
SQL> CREATE UNIQUE INDEX ssb.test_idx ON ssb.test(ENAME, EADD1, EADD2, EADD3, EADD4, CITY) COMPRESS ADVANCED HIGH;
```

```
SQL> select owner,segment_name,segment_type,bytes from dba_segments where owner='SSB' and segment_name in ('TEST','TEST_IDX');
```

OWNER	SEGMENT_NAME	SEGMENT_TYPE	BYTES
SSB	TEST	TABLE	6291456
SSB	TEST_IDX	INDEX	196608

```
SQL> drop index ssb.test_idx;
```

```
SQL> exit;
```

Advanced Compression: Oracle SecureFiles



La funcionalidad SecureFiles es un rediseño completo de la implementación del almacenamiento de objetos grandes (LOB) en Oracle 12c. El almacenamiento LOB original, conocido como BASICFILE, sigue siendo el método de almacenamiento predeterminado en versión 11g, pero la palabra clave SECUREFILE habilita el nuevo método de almacenamiento, que permite el cifrado y ahorro de espacio mediante compresión y deduplicación. En 12c es posible establecer a nivel de parámetros de base de datos el poder manejar los campos tipo LOB (Securefiles) como predeterminado en la creación de cualquier tipo LOB. En 18c y 19c por defecto se crean como SECUREFILES.

Para utilizar los campos tipo Securefiles los tablespaces tienen que ser definidos como SEGMENT SPACE MANAGEMENT AUTO;

Para el siguiente ejemplo se crean dos tablespaces (uno para cargar ficheros en base de datos BASICFILES y otro como SECUREFILES).

Se cargarán unos cuantos ficheros de texto .TXT existentes en la carpeta /home/oracle/sf/docs en LOB (BASICFILES) y posteriormente los insertaremos en tipo LOB SECUREFILES (comprimidos).

```
[oracle@single19c ~]$ sqlplus / as sysdba

SQL*Plus: Release 19.0.0.0.0 - Production on Fri Sep 4 12:50:16 2020
Version 19.8.0.0.0

Copyright (c) 1982, 2020, Oracle. All rights reserved.

Connected to:
Oracle Database 19c EE Extreme Perf Release 19.0.0.0.0 - Production
Version 19.8.0.0.0

SQL> alter session set container=SOE;

Session altered.

SQL> show pdbs

  CON_ID CON_NAME                                OPEN MODE  RESTRICTED
  -----
  5 SOE                                           READ WRITE NO
```

```
SQL> CREATE TABLESPACE securefiles
  DATAFILE '+DATA/'
  SIZE 300M REUSE
  EXTENT MANAGEMENT LOCAL
  UNIFORM SIZE 4M
  SEGMENT SPACE MANAGEMENT AUTO;

Tablespace created.
```



```
SQL> CREATE TABLESPACE basicfiles
      DATAFILE '+DATA/'
      SIZE 300M REUSE
      EXTENT MANAGEMENT LOCAL
      UNIFORM SIZE 4M
      SEGMENT SPACE MANAGEMENT AUTO;
```

Tablespace created.

-- Crear un usuario para la carga de los ficheros de texto .txt

```
SQL> CREATE USER sf
      IDENTIFIED BY We1c0m3_We1c0m3_
      DEFAULT TABLESPACE sysaux
      TEMPORARY TABLESPACE temp
      QUOTA UNLIMITED ON sysaux
      QUOTA UNLIMITED ON basicfiles
      QUOTA UNLIMITED ON securefiles
/
```

User created

```
SQL> GRANT dba TO sf;
```

```
SQL> GRANT EXECUTE ANY PROCEDURE, CREATE ANY DIRECTORY TO sf;
```

Se crea con el usuario sf un directorio para almacenar los ficheros que luego se utilizarán para cargarlos en la base de datos tipo Securefiles.

-- Crear el directorio de objetos donde estarán los ficheros de texto .txt

```
SQL> conn sf/We1c0m3_We1c0m3_@soe;
```

```
SQL> CREATE OR REPLACE DIRECTORY sf_docs
      AS '/home/oracle/sf/docs';
```

Directory created.

Creamos la tabla TICKETS que contendrá los campos LOB tipo BASICFILES y dos tablas que contendrán los campos LOB tipo SECUREFILES: SECURE_TICKETS (no comprimido) y SECURE_TICKETS_COMP (comprimido).

```
SQL> CREATE TABLE sf.tickets (
      tkt_id          NUMBER
      ,description    VARCHAR2(30)
      ,submit_dtm     TIMESTAMP
      ,status         VARCHAR2(8)
      ,document       BLOB
)
LOB(document) STORE AS BASICFILE (TABLESPACE basicfiles)
```



```

;
Table created
SQL> CREATE TABLE sf.secure_tickets (
    tkt_id          NUMBER
    ,description    VARCHAR2(30)
    ,submit_dtm     TIMESTAMP
    ,status         VARCHAR2(8)
    ,document       BLOB
)
  LOB(document)
    STORE AS SECUREFILE (
      TABLESPACE securefiles
    )
;

```

Table created.

```

SQL> CREATE TABLE sf.secure_tickets_comp (
    tkt_id          NUMBER
    ,description    VARCHAR2(30)
    ,submit_dtm     TIMESTAMP
    ,status         VARCHAR2(8)
    ,document       BLOB
)
  LOB(document)
    STORE AS SECUREFILE (
      TABLESPACE securefiles
      COMPRESS HIGH
    )
;

```

Table created.

Utilizaremos un procedimiento para cargar los ficheros de texto .TXT existentes en la carpeta /home/oracle/sf/docs en la tabla TICKETS en tipo BASICFILES.

```

SQL> CREATE OR REPLACE PACKAGE sf.pkg_securefiles
AS
    PROCEDURE AddTroubleTicket (
        tkt_id          IN sf.tickets.tkt_id%TYPE
        ,description    IN sf.tickets.description%TYPE
        ,submit_dts     IN VARCHAR2
        ,status         IN sf.tickets.status%TYPE
        ,docFileName    IN VARCHAR2
    );
END pkg_securefiles;
/

```



Package created.

```
SQL> CREATE OR REPLACE PACKAGE BODY sf.pkg_securefiles
AS
```

```
    PROCEDURE LoadBFILEIntoLOB (
        src_dir    IN      VARCHAR2
        ,src_file   IN      VARCHAR2
        ,target_lob IN OUT  BLOB
    )
```

```
    IS
```

```
        src_loc    BFILE    := BFILENAME(src_dir, src_file);
        load_amt    INTEGER := 4000;
```

```
    BEGIN
```

```
        -- Open the source document file in read-only mode
        DBMS_LOB.OPEN(
            file_loc => src_loc
            ,open_mode => DBMS_LOB.LOB_READONLY
        );
```

```
        -- Calculate the size of the external BFILE
        load_amt := DBMS_LOB.GETLENGTH(file_loc => src_loc);
```

```
        -- Load the LOB from the source file
        DBMS_LOB.LOADFROMFILE(target_lob, src_loc, load_amt);
```

```
        -- Close the opened BFILE external LOB
        DBMS_LOB.FILECLOSE(file_loc => src_loc);
```

```
    EXCEPTION
```

```
        WHEN OTHERS THEN
```

```
            DBMS_OUTPUT.PUT_LINE('LoadLOBFromFile Error: ' || SQLCODE || ' - '
|| SQLERRM);
```

```
    END LoadBFILEIntoLob;
```

```
    PROCEDURE AddTroubleTicket (
```

```
        tkt_id      IN sf.tickets.tkt_id%TYPE
        ,description IN sf.tickets.description%TYPE
        ,submit_dts   IN VARCHAR2
        ,status       IN sf.tickets.status%TYPE
        ,docFileName  IN VARCHAR2
    )
```

```
    )
```

```
    IS
```

```
        submit_dtm  TIMESTAMP;
        docBlob      BLOB;
```

```
    BEGIN
```

```
        -- Calculate timestamp value
        submit_dtm := TO_TIMESTAMP(submit_dts, 'yyyy-mm-dd hh24:mi:ss');
```

```
        -- Add new row, returning references to the document and image BLOBs
        INSERT INTO sf.tickets
```



```

VALUES (tkt_id, description, submit_dtm, status, EMPTY_BLOB())
RETURNING document INTO docBlob;

-- Build the document LOB from the supplied file name
LoadBFILEIntoLOB('SF_DOCS', docFileName, docBlob);

EXCEPTION
  WHEN OTHERS THEN
    DBMS_OUTPUT.PUT_LINE('Severe error! ' || SQLCODE || ' - ' ||
SQLERRM);

    END AddTroubleTicket;
END pkg_securefiles;
/

Package body created.

```

Cargamos los ficheros de texto .TXT existentes en la carpeta /home/oracle/sf/docs en la tabla TICKETS, para posteriormente migrarlos a formato Securefiles.

```

SQL> SET SERVEROUTPUT ON

SQL> BEGIN

  sf.pkg_securefiles.AddTroubleTicket (
    tkt_id => 101
    ,description => 'Trouble Ticket 101'
    ,submit_dts => '2008-12-31 23:45:00'
    ,status => 'OPEN'
    ,docFileName => 'prueba.txt'
  );

  sf.pkg_securefiles.AddTroubleTicket (
    tkt_id => 102
    ,description => 'Trouble Ticket 102'
    ,submit_dts => '2009-01-04 00:00:00'
    ,status => 'OPEN'
    ,docFileName => 'prueba1.txt'
  );

  sf.pkg_securefiles.AddTroubleTicket (
    tkt_id => 103
    ,description => 'Trouble Ticket 103'
    ,submit_dts => '2009-01-02 00:00:00'
    ,status => 'OPEN'
    ,docFileName => 'prueba2.txt'
  );

  sf.pkg_securefiles.AddTroubleTicket (
    tkt_id => 104
    ,description => 'Trouble Ticket 104'
    ,submit_dts => '2009-01-14 12:30:00'
  );

```



```

        ,status => 'OPEN'
        ,docFileName => 'prueba3.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 105
        ,description => 'Trouble Ticket 105'
        ,submit_dts => '2009-01-09 00:00:00'
        ,status => 'OPEN'
        ,docFileName => 'prueba4.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 106
        ,description => 'Trouble Ticket 106'
        ,submit_dts => '2009-01-11 00:00:00'
        ,status => 'OPEN'
        ,docFileName => 'prueba5.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 107
        ,description => 'Trouble Ticket 107'
        ,submit_dts => '2009-01-16 00:00:00'
        ,status => 'OPEN'
        ,docFileName => 'prueba6.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 108
        ,description => 'Trouble Ticket 108'
        ,submit_dts => '2009-01-12 00:00:00'
        ,status => 'OPEN'
        ,docFileName => 'prueba7.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 109
        ,description => 'Trouble Ticket 109'
        ,submit_dts => '2009-01-02 00:00:00'
        ,status => 'OPEN'
        ,docFileName => 'prueba8.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 110
        ,description => 'Trouble Ticket 110'
        ,submit_dts => '2009-01-14 12:45:00'
        ,status => 'OPEN'
        ,docFileName => 'prueba9.txt'
    );

    sf.pkg_securefiles.AddTroubleTicket (
        tkt_id => 111
        ,description => 'Trouble Ticket 111'
        ,submit_dts => '2009-01-14 12:45:00'
    );

```



```

        ,status => 'OPEN'
        ,docFileName => 'prueba10.txt'
    );
    COMMIT;
END;
/

```

PL/SQL procedure successfully completed.

Analizamos el schema SF mediante DBMS_STATS.GATHER_SCHEMA_STATS.

```

SQL> BEGIN
    DBMS_STATS.GATHER_SCHEMA_STATS(ownname => 'sf', CASCADE => TRUE);
END;
/

```

PL/SQL procedure successfully completed.

Insertamos en las tablas creadas como SECUREFILES a partir de los datos en la tabla creada como BASICFILES mediante insert ... select ...

```

SQL> INSERT INTO sf.secure_tickets
SELECT * FROM sf.tickets;

11 rows created.

SQL> INSERT INTO sf.secure_tickets_comp
SELECT * FROM sf.tickets;

11 rows created.

SQL> COMMIT;

```

Comprobar los metadatos creados en ambas tablas accediendo a la vista DBA_SEGMENTS:

```

SET PAGESIZE 1000
SET LINESIZE 140
set serveroutput on
-----
-- View: DBA_SEGMENTS
-- Shows metadata about individual BasicFile and SecureFile segments
-----

TTITLE 'LOB Segment Information|(from DBA_SEGMENTS)'
COL segment_name          FORMAT A30          HEADING 'Segment Name'
COL segment_type          FORMAT A20          HEADING 'Segment|Type'
COL segment_subtype       FORMAT A20          HEADING 'Segment|SubType'
COL partition_name        FORMAT A12          HEADING 'Partition|Name'
COL tablespace_name       FORMAT A12          HEADING 'Tablespace'
SELECT

```




```

        segment_name
        ,segment_type
        ,segment_subtype
        ,partition_name
        ,tablespace_name
    FROM dba_segments
    WHERE owner = 'SF'
    ORDER BY segment_name
;
TTITLE OFF

TTITLE OFF
   3   4   5   6   7   8   9  10
Wed Jan 12
page      1

LOB Segment
Information
                                (from DBA_SEGMENTS)

Partition
Segment Name                    Segment
Tablespace                     Type          Segment
                                Name
-----
SECURE_TICKETS                  TABLE      ASSM
SYSAUX
SECURE_TICKETS_COMP            TABLE      ASSM
SYSAUX
SYS_IL0000075331C00005$$      LOBINDEX    ASSM
BASICFILES
SYS_IL0000075334C00005$$      LOBINDEX    ASSM
SECUREFILES
SYS_IL0000075337C00005$$      LOBINDEX    ASSM
SECUREFILES
SYS_LOB0000075331C00005$$      LOBSEGMENT  ASSM
BASICFILES
SYS_LOB0000075334C00005$$      LOBSEGMENT  SECUREFILE
SECUREFILES
SYS_LOB0000075337C00005$$      LOBSEGMENT  SECUREFILE
SECUREFILES
TICKETS                        TABLE      ASSM
SYSAUX

9 rows selected.

```

Basado en la información obtenida de la vista DBA_SEGMENTS modificar el nombre de los segmentos (LOB) correspondientes al tablespace SECUREFILES en el bloque PL/SQL más abajo, concretamente la parte resaltada en color rojo. A continuación, se muestra una salida de ejemplo:



SYS_LOB0000075334C00005\$\$	LOBSEGMENT	SECUREFILE
SECUREFILES		
SYS_LOB0000075337C00005\$\$	LOBSEGMENT	SECUREFILE
SECUREFILES		

Antes de ejecutar el siguiente bloque, modificar el código PL/SQL y añadir los valores recuperados en la consulta anterior reemplazando el texto en rojo.

```
SQL> declare
    segment_size_block NUMBER;
    segment_size_byte NUMBER;
    used_block NUMBER;
    used_byte NUMBER;
    expired_block NUMBER;
    expired_byte NUMBER;
    unexpired_block NUMBER;
    unexpired_byte NUMBER;
begin
    dbms_space.space_usage ('SF', 'SYS_LOB0000075334C00005$$', 'LOB',
    dbms_space.spaceusage_exact, segment_size_block,
    segment_size_byte, used_block, used_byte, expired_block, expired_byte,
    unexpired_block, unexpired_byte, null);
    dbms_output.put_line('segment_size_blocks = '||segment_size_block);
    dbms_output.put_line('segment_size_bytes = '||segment_size_byte);
    dbms_output.put_line('used_blocks = '||used_block);
    dbms_output.put_line('used_bytes = '||used_byte);
    dbms_output.put_line('expired_blocks = '||expired_block);
    dbms_output.put_line('expired_bytes = '||expired_byte);
    dbms_output.put_line('unexpired_blocks = '||unexpired_block);
    dbms_output.put_line('unexpired_bytes = '||unexpired_byte);
end;
/

segment_size_blocks = 3072
segment_size_bytes = 25165824
used_blocks = 2468
used_bytes = 20217856
expired_blocks = 604
expired_bytes = 4947968
unexpired_blocks = 0
unexpired_bytes = 0

PL/SQL procedure successfully completed.

segment_size_blocks = 512
segment_size_bytes = 4194304
used_blocks = 87
used_bytes = 712704
expired_blocks = 425
expired_bytes = 3481600
unexpired_blocks = 0
```



```
unexpired_bytes = 0
```

```
PL/SQL procedure successfully completed.
```

Observar la ocupación de los segmentos de la columna 'segment_size_bytes' de ambos objetos por el nombre de los segmentos LOBS.

Consultas sobre tablas Comprimidas/No Comprimidas

En el siguiente ejemplo se ejecutarán varias sentencias sobre tablas Comprimidas y No comprimidas para ver el rendimiento de ambas. ACO nos proporciona beneficios tanto de reducción del almacenamiento, como rendimiento en consultas al tener que leer menos bloques en memoria de los Buffers de la SGA.

Con las siguientes sentencias, vemos que hay dos tablas (LINEORDER y LINEORDER_NO_ACO) una comprimida con compresión BASIC y otra sin compresión. De igual forma vemos la ocupación de bloques de cada tabla.

```
$ sqlplus ssb/ssb@soe;

SQL*Plus: Release 19.0.0.0.0 - Production on Wed Jan 12 16:07:29 2022
Version 19.13.0.0.0

Copyright (c) 1982, 2021, Oracle. All rights reserved.

Last Successful login time: Wed Jan 12 2022 13:13:03 +00:00

Connected to:
Oracle Database 19c EE Extreme Perf Release 19.0.0.0.0 - Production
Version 19.13.0.0.0

SQL> set linesize 1000
SQL> col table_name format a30

SQL> SELECT table_name,compression, compress_for
FROM   user_tables
WHERE  table_name like 'LINEORDE%';

TABLE_NAME                                COMPRESS COMPRESS_FOR
-----
LINEORDER                                ENABLED  BASIC
LINEORDER_ACO                            ENABLED  ADVANCED
LINEORDER_NO_ACO                         DISABLED

SQL> SELECT table_name,
             compression,
             num_rows,
             blocks,
```



```

        empty_blocks
FROM    user_tables
WHERE   table_name LIKE ('LINEORDER%')
ORDER BY 1;

```

TABLE_NAME	COMPRESS	NUM_ROWS	BLOCKS	EMPTY_BLOCKS
LINEORDER	ENABLED	29999800	223987	0
LINEORDER_ACO	ENABLED	1000000	8416	0
LINEORDER_NO_ACO	DISABLED	29999870	417832	0

Ejecutar las siguientes sentencias para ver como responden sobre tablas comprimidas/no comprimidas.

```

SQL> set timing on

SQL> select /* no compresion */ max(LO_ORDTOTALPRICE) most_expensive_order From
LINEORDER_NO_ACO where LO_PARTKEY=300023;

MOST_EXPENSIVE_ORDER
-----
          40614642

Elapsed: 00:00:29.68

SQL> select /* compresion */ max(LO_ORDTOTALPRICE) most_expensive_order From
LINEORDER where LO_PARTKEY=300023;

MOST_EXPENSIVE_ORDER
-----
          40614642

Elapsed: 00:00:00.17

```

Otras dos consultas un poco más complejas

```

SQL> select /* no compresion */ d_date,sum(l.lo_revenue) "Total Revenue"
From    LINEORDER_NO_ACO l, DATE_DIM d
Where   l.lo_orderdate = d.d_datekey
and     D_DAYNUMINMONTH = 25
and     d.d_month = 'December'
group by d_date
order by d_date;

D_DATE                Total Revenue
-----
December 25, 1992     4.6038E+10
December 25, 1993     4.3746E+10
December 25, 1994     4.5026E+10

```



```
December 25, 1995      4.5589E+10
December 25, 1996      4.6722E+10
December 25, 1997      4.4994E+10
```

6 rows selected.

Elapsed: 00:00:29.77

```
SQL> select /* compresion */ d_date,sum(l.lo_revenue) "Total Revenue"
From   LINEORDER l, DATE_DIM d
Where  l.lo_orderdate = d.d_datekey
and    D_DAYNUMINMONTH = 25
and    d.d_month = 'December'
group by d_date
order by d_date;
```

D_DATE	Total Revenue
December 25, 1992	4.6038E+10
December 25, 1993	4.3746E+10
December 25, 1994	4.5026E+10
December 25, 1995	4.5589E+10
December 25, 1996	4.6722E+10
December 25, 1997	4.4994E+10

6 rows selected.

Elapsed: 00:00:02.84

Proceso de compresión de expdp

Desde la versión 11g, Data Pump permite comprimir el backup antes de escribir a fichero dump con el parámetro 'compression'.

El parámetro 'compression' puede contener 4 valores:

- ALL
- DATA_ONLY
- METAGATA_ONLY
- NONE

Utilizando el valor 'ALL', el tamaño del fichero de backup puede ser reducido hasta 10 veces. El tiempo de expdp se incrementará con respecto al export sin compresión.

Vamos a crear dos ficheros de parámetros para EXPDP, uno con compresión y otro si ella. Para crear estos ficheros use el editor vi o cree el fichero con un editor de texto local y transfíralo a la máquina con SFTP.

El contenido de exp_pdbsoe_comp.par (con compresión) debe ser el siguiente.

```
directory=DATA_PUMP_DIR
```



```
dumpfile=EXP_PDBSOE_COMP%U.dmp
logfile=EXP_PDBSOE_COMP.log
COMPRESSION=ALL
COMPRESSION_ALGORITHM=HIGH
SCHEMAS=SOE
Exclude=materialized_view
```

El contenido de exp_pdbsoe_nocomp.par (sin compresión) debe ser el siguiente.

```
directory=DATA_PUMP_DIR
dumpfile=EXP_PDBSOE_NOCOMP%U.dmp
logfile=EXP_PDBSOE_NOCOMP.log
SCHEMAS=SOE
Exclude=materialized_view
```

Por defecto existe un directorio con nombre 'DATA_PUMP_DIR'. Ejecutar la siguiente sentencia para ver el PATH completo de disco.

```
col owner format a20;
col directory_name format a30;
col directory_path format a65;

SQL> select * from dba_directories where DIRECTORY_NAME= 'DATA_PUMP_DIR';
```

OWNER	DIRECTORY_NAME	DIRECTORY_PATH	ORIGIN_CON_ID
SYS	DATA_PUMP_DIR	/u01/app/oracle/product/19.0.0.0/dbhome_1/rdbms/log/	1

Lanzar los comandos expdp con los ficheros .PAR creado en el paso anterior, para la pluggable database SOE.

```
$ expdp system/we1c0m3_we1c0m3_@SOE parfile=exp_pdbsoe_comp.par
$ expdp system/we1c0m3_we1c0m3_@SOE parfile=exp_pdbsoe_nocomp.par
```

Comprobar el tiempo que tarda cada uno y el espacio ocupado por el fichero .dmp generado.

```
31060 -rw-r----- 1 oracle asmadmin 31801344 Jan 12 16:21
EXP_PDBSOE_COMP01.dmp
87156 -rw-r----- 1 oracle asmadmin 89243648 Jan 12 16:23
EXP_PDBSOE_NOCOMP01.dmp
```

