



ORACLE

# Oracle Database MultiTenant/MultiModel/InMemory/ACO Workshop

Juan Carlos Díaz

José Vázquez

Francisco Alvarez

Francisco Rivas

Madrid 18 de Enero de 2022

## Safe harbor statement

---

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions.

The development, release, timing, and pricing of any features or functionality described for Oracle's products may change and remains at the sole discretion of Oracle Corporation.

# Overview

- Multitenant
- MultiModel - JSON
- InMemory
- ACO

# Overview

- Multitenant
- Multimodel - JSON
- InMemory



Oracle 12.1

Oracle 12.2

Oracle 18

Oracle 19

### Deprecation of Non-CDB Architecture

The non-CDB architecture was deprecated in Oracle Database 12c. It can be desupported and unavailable in a release after Oracle Database 19c .

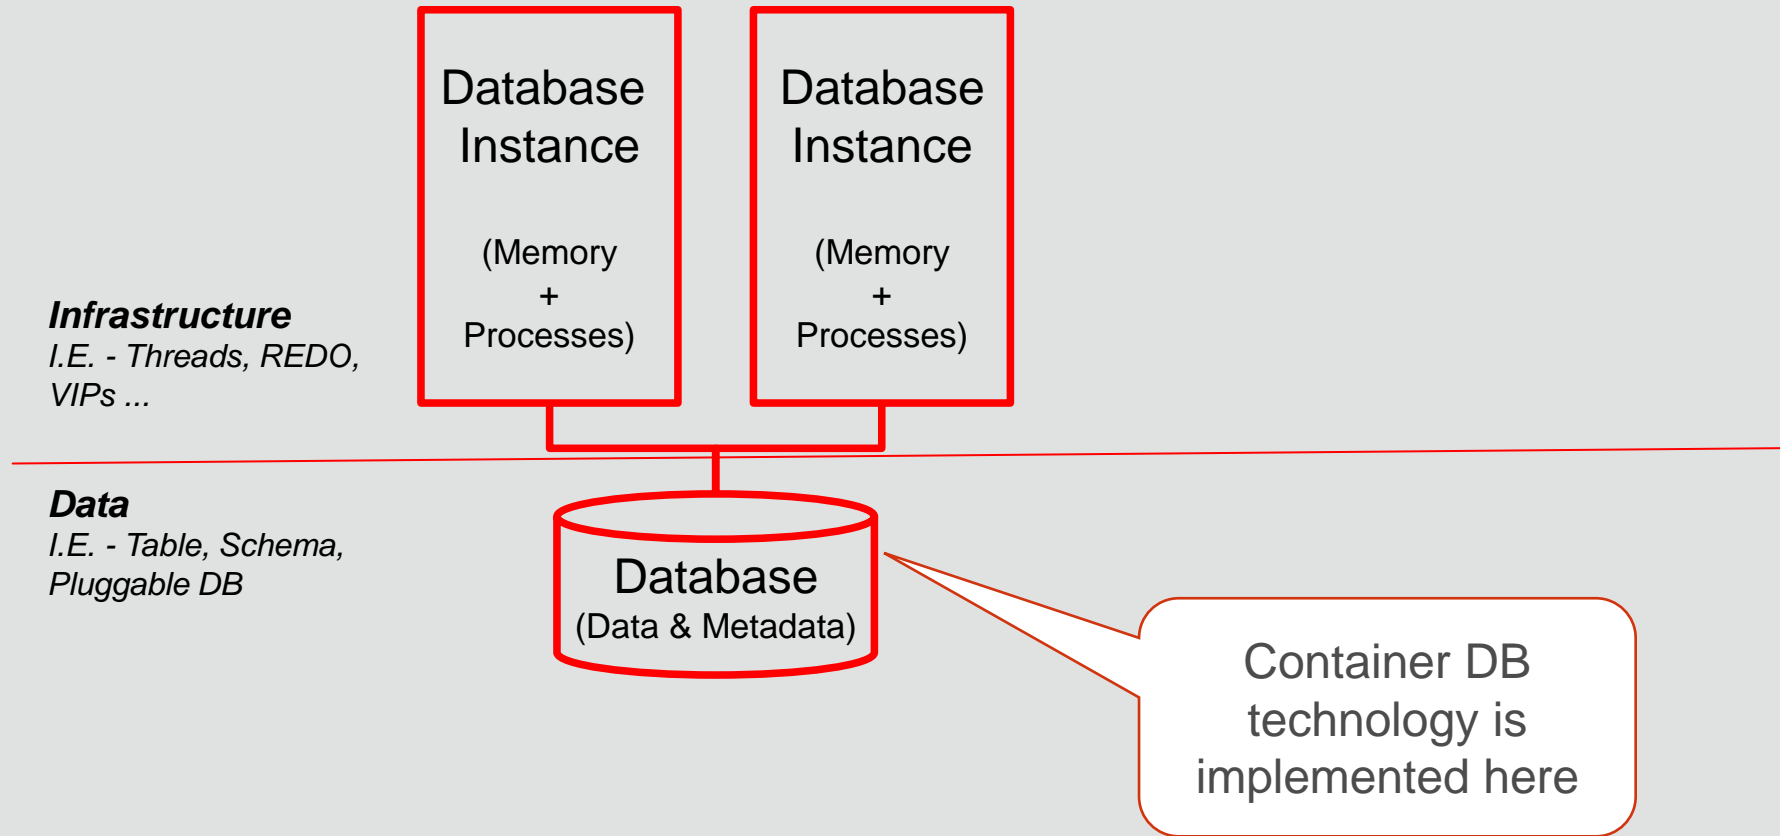
***“Oracle recommends use of the CDB architecture”***



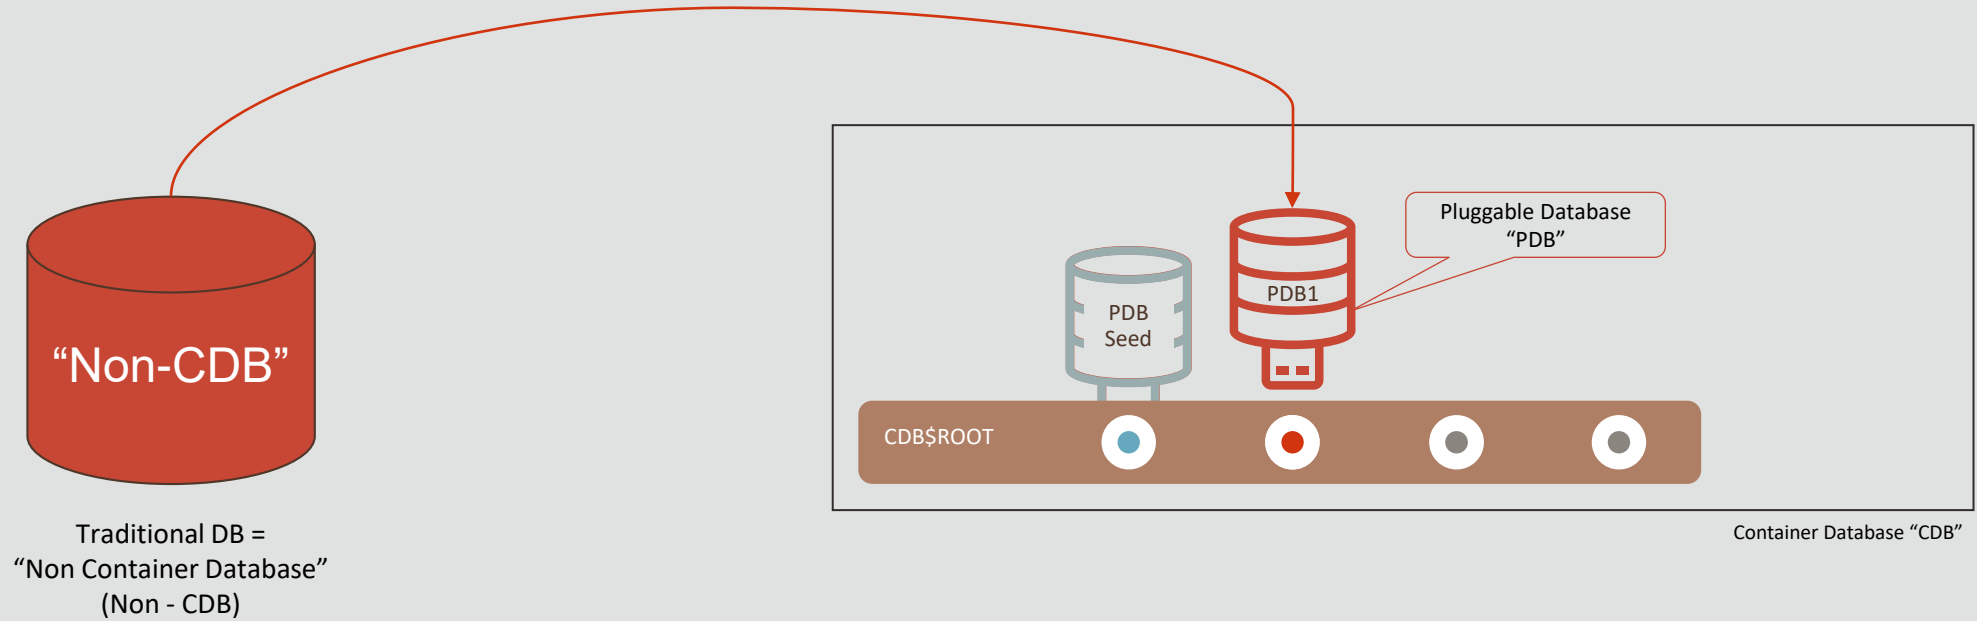
19 - Sept - 2019

*“The Oracle Database non-CDB architecture  
will be **de-supported** from Oracle Database 20c onwards”*

# CDB architecture

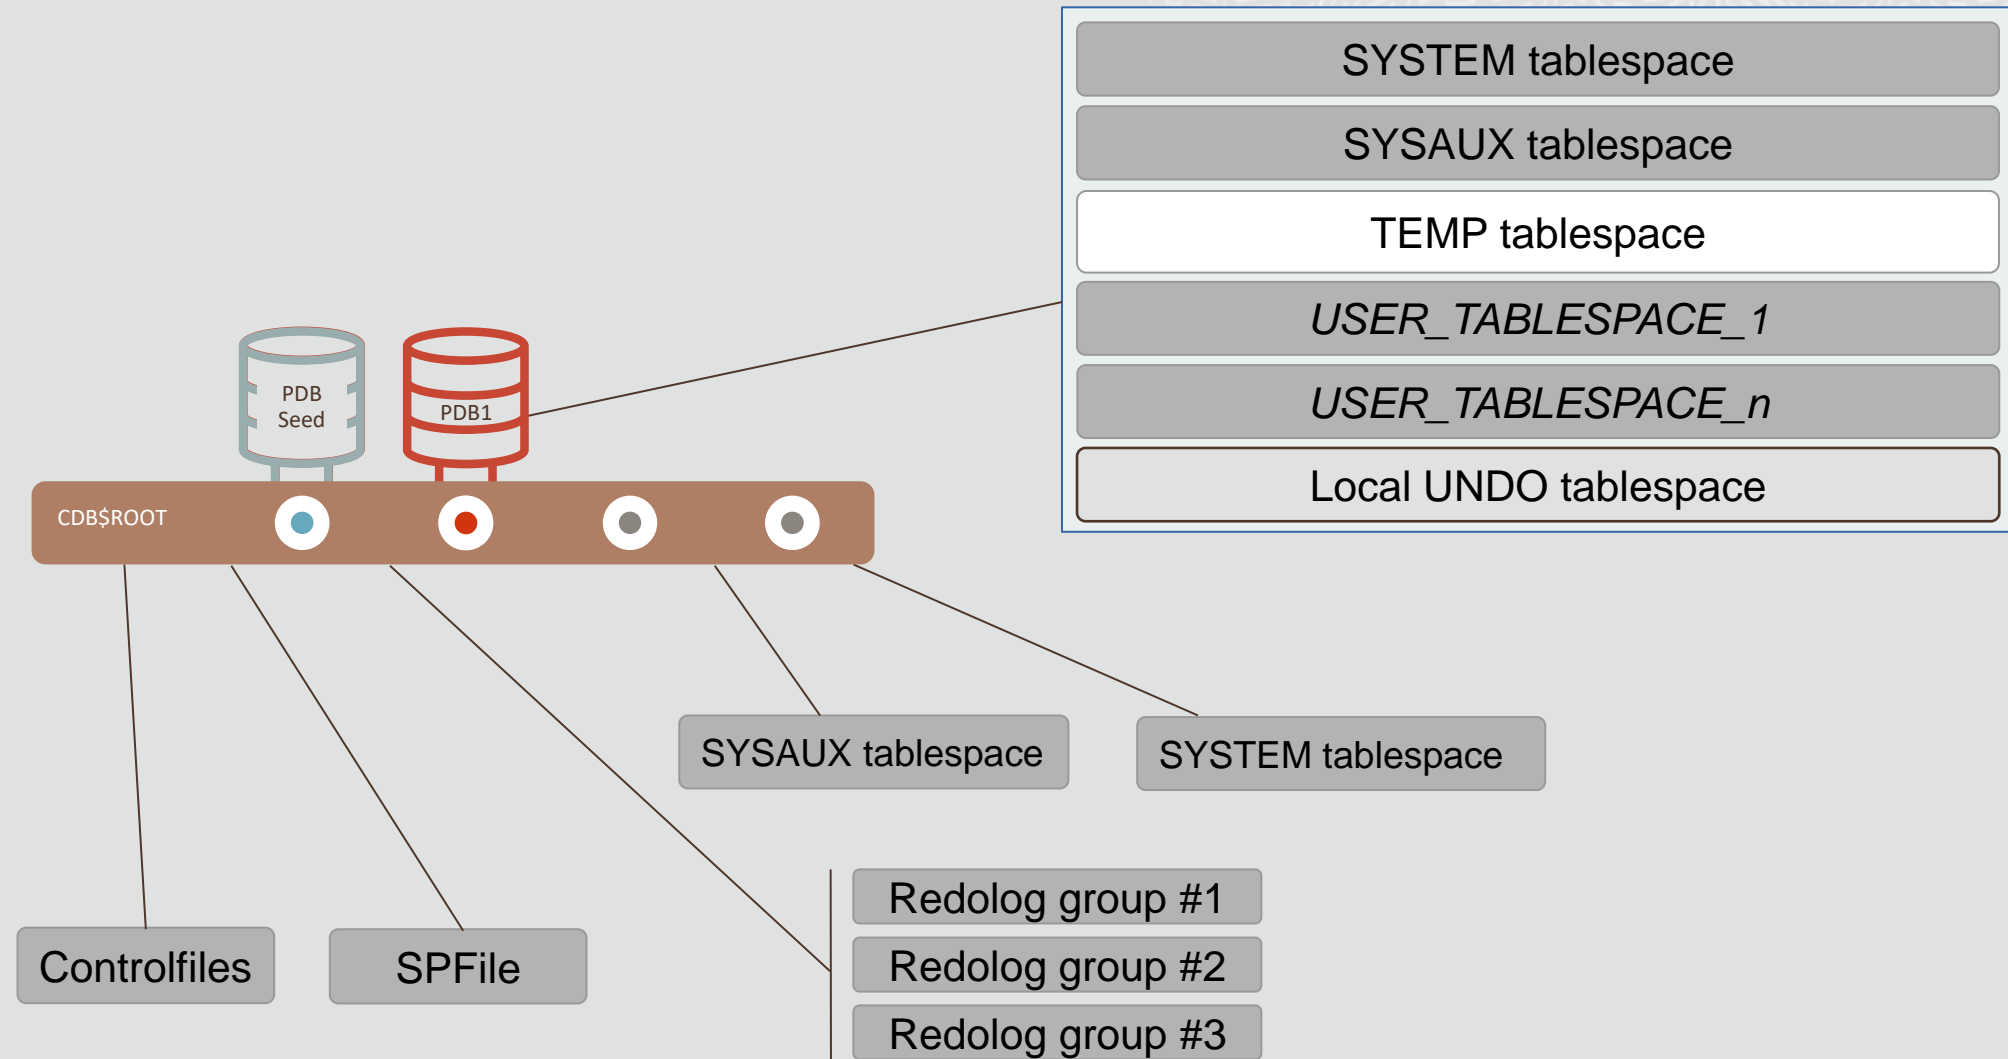


# From Non-CDB to CDB architecture



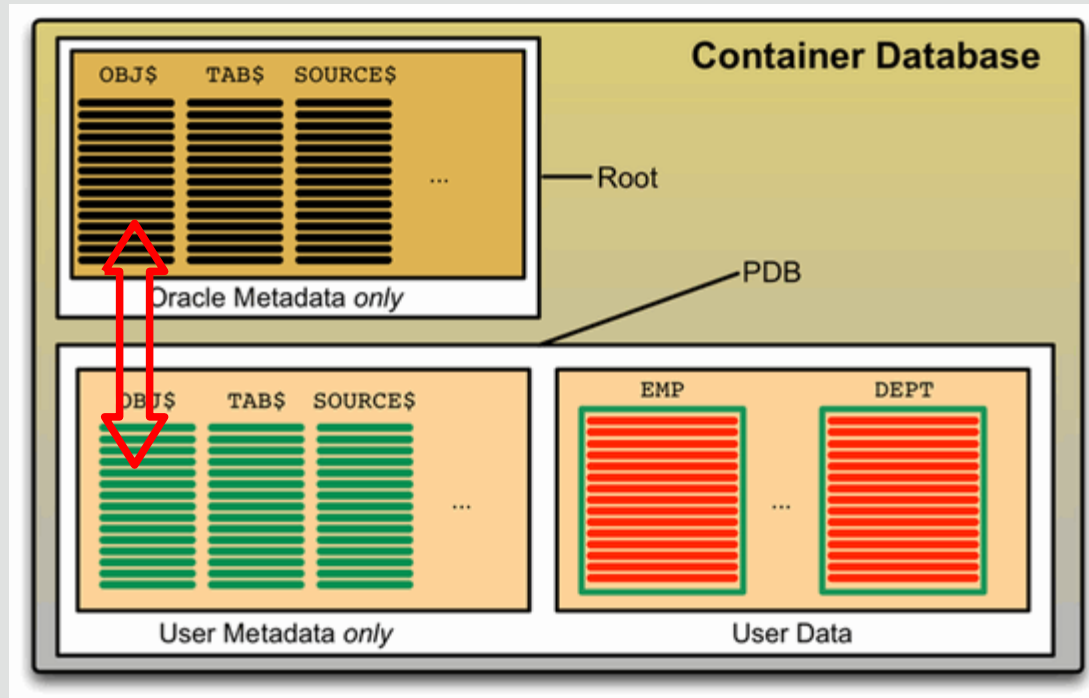
- ✓ 100% Logically Isolated.
- ✓ Same DB as it was. Same identity, same behavior.

# CDB architecture





## How multiple dictionaries?














Root dictionary and PDBs dictionaries are :

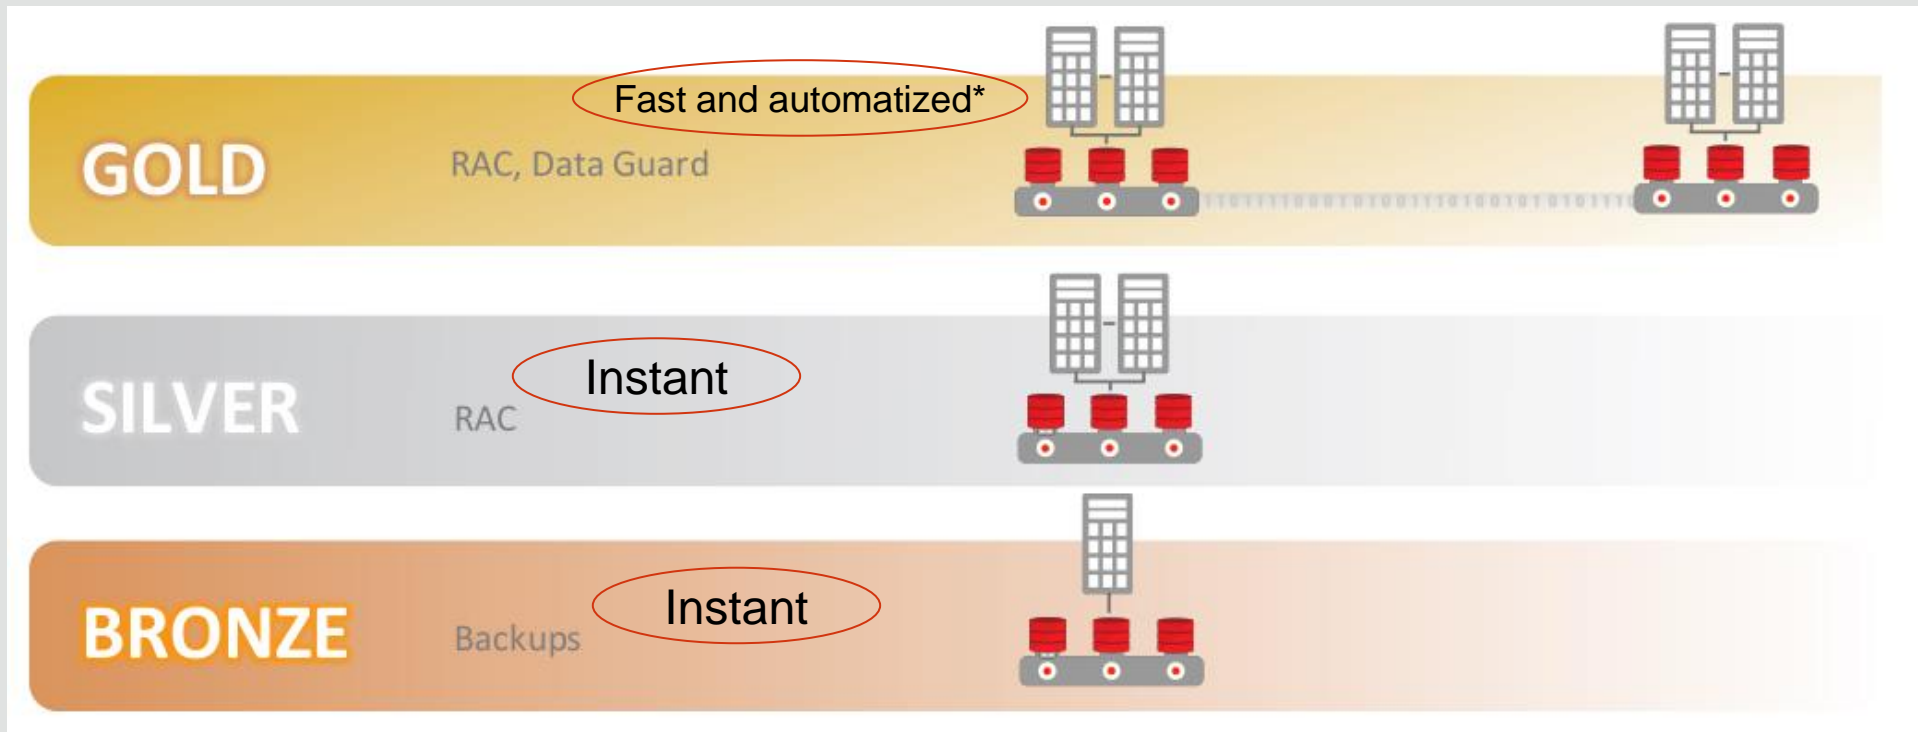
- ✓ Logically independent
- ✓ Internally connected with stubs for common data
- ✓ Same structures on each dictionary
- ✓ Different rows on each dictionary, thus “horizontal partitioned”.

## PDB Means mobility

Multiple SYSTEM & SYSAUX tablespaces  
Where is What?

Feature	Stored @Root		Stored @PDB	
ASH*				
AWR		Root perf. info		PDB perf. Info
ADDM*				
Segment Advisor				
Optimizer statistics				
DBReplay and/or SPA				

## PDB Means mobility




# Main Use cases

1. More power under the same hardware (multitenant)
2. Resources control (multitenant & singletenant)
3. Faster, easier and new maintenances (multitenant & singletenant)

# Resource isolation with the deprecated architecture


Non-CBD  
19c - 30 DB




- Competing use of CPUs (full cpu per instance) or underperforming use of CPU (caging)
- Unused & unshared SGA
- Unused & unshared PGA
- $N$  sets of background processes

# More power under the same hardware

Non-CBD  
19c - 30 DB



- Competing use of CPUs (full cpu per instance) or underperforming use of CPU (caging)
- Unused & unshared SGA
- Unused & unshared PGA
- *N* sets of background processes



- ✓ Online and immediate CPU rebalance
- ✓ Online and immediate SGA rebalance
- ✓ Shared PGA free space
- ✓ 1 set of background processes



# More power under the same hardware



## Unconsolidated


- Overheads replicated for each workload
- Competing workloads interfere with each other
- Net throughput is limited and gridlock is possible



## PDB Consolidation

- Single, shared set of overheads
- Resource sharing vs isolation
- Coordination and cooperation vs competition

# More power under the same hardware







# CPU Management

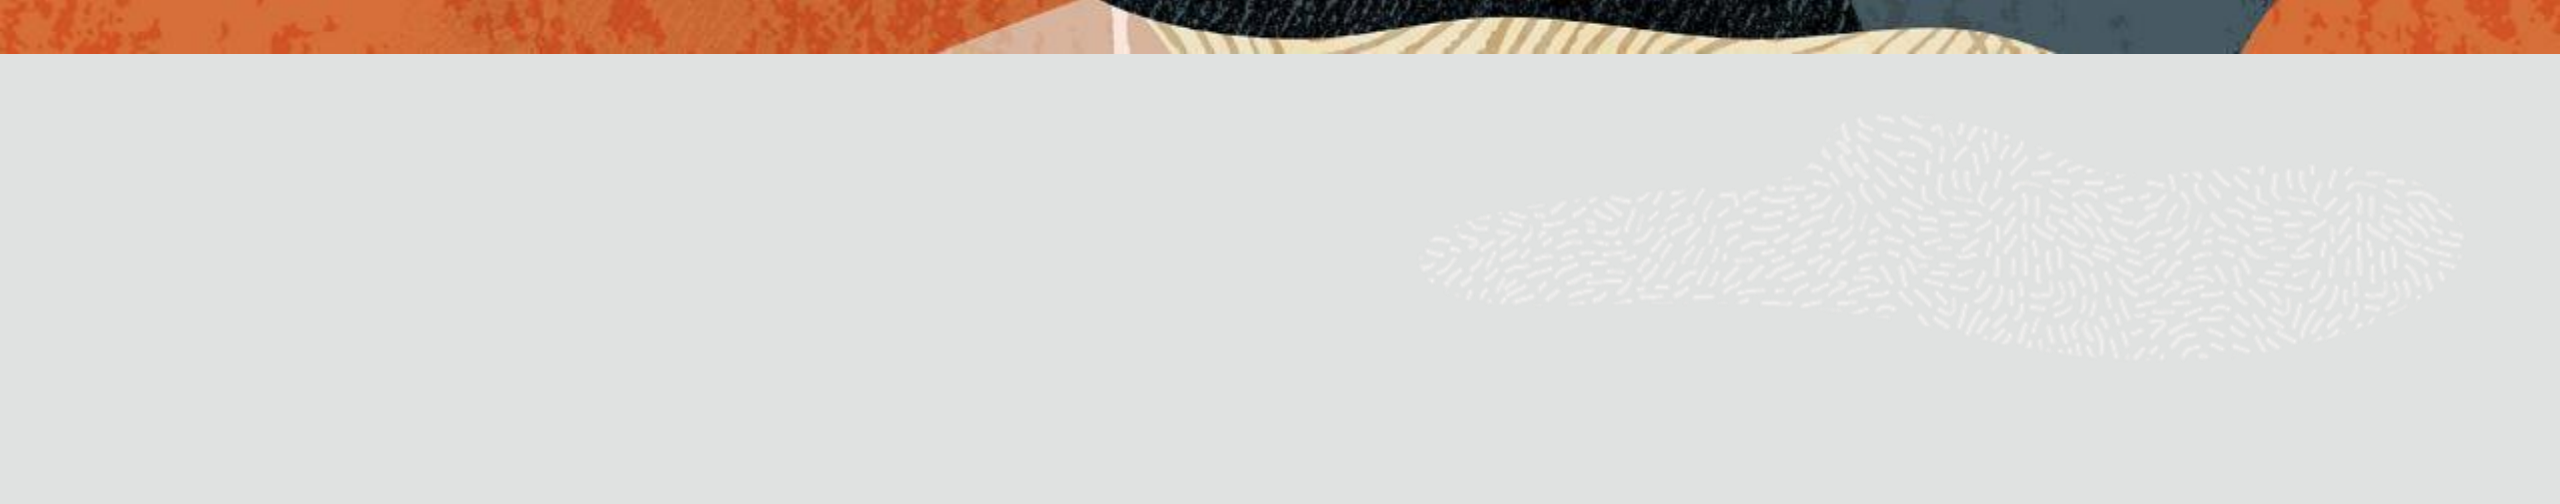
## Why PDB online resource sharing is so capable

1. Datawarehouse reports executed at specific time-frames.
2. Reports in time for the CEO = buy HW for peak loads
3. CPU and memory wasted during idle time.



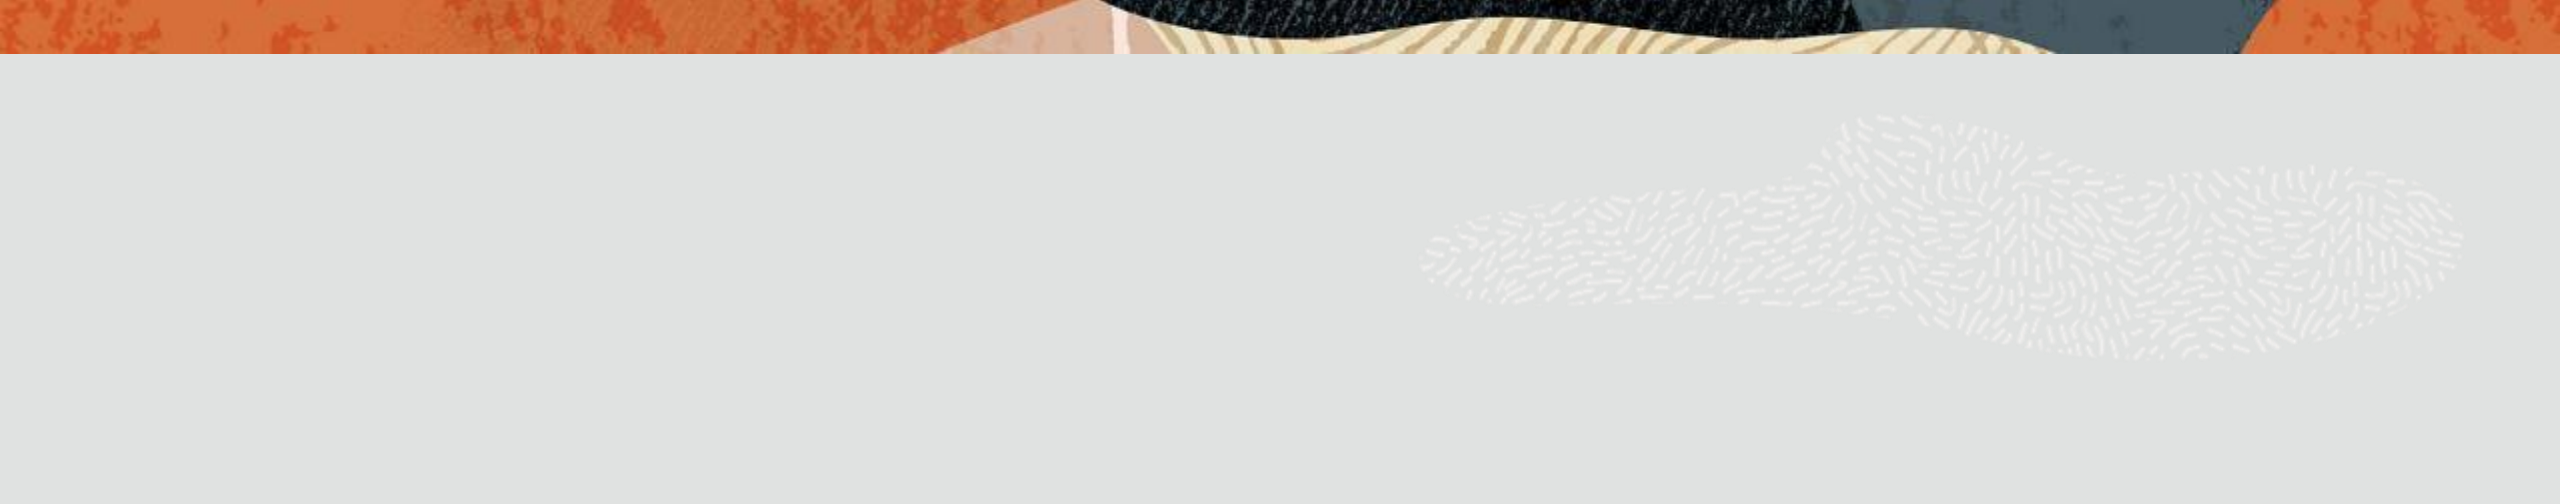
# CPU Management : Greater flexibility as consolidation increases



- 
1. More power under the same hardware (multitenant)
  2. **Resources control** (multitenant & singletenant)
  3. Faster, easier and new maintenances (multitenant & singletenant)

# Real online control of your resources

- ✓ **CPU** : No more fixed CPU values per database; defined dynamically by priorities. Can set minimums. Can set maximums. Modifiable online. Modifiable by a schedule.
- ✓ **Memory** : All the memory accessible by all databases. SGA and PGA can be set with minimums, maximums and/or fixed values.
- ✓ **I/O** : Can define IO rate limits (both IOPS and/or MBPS) for each PDB. Can be dynamically set though priorities with Exadata (*I.O.R.M at PDB level*)
- ✓ **Parallel Execution Servers** : Parallelism capacity can be controled dynamically between PDBs by priorities.

- 
1. More power under the same hardware (multitenant)
  2. Power *with* control (multitenant & singletenant)
  3. **Faster, easier and *new* maintenances** (multitenant & singletenant)

# New features brought by CDB architecture

- ✓ **Fast database provisioning** : ~1 minute vs 45 minutes with old architecture
- ✓ **Online PDB relocation** : 2 simple SQL commands. Powerful migration option.
- ✓ **Online PDB clone (local or remote)**: 1 or 2 simple SQL commands. RMAN not needed
- ✓ **PDB Snapshot clone**: 1 simple SQL command. Immediate. Space saver.
- ✓ **ASM Split Mirror**: 2 simple SQL commands. Immediate read/write clone.
- ✓ **Many B&R features** (PDB level recovery, PDB level flashback, etc)

# Workshop

MultiTenant + JSON + InMemory + ACO

<https://github.com/OracleDataManagementSpain/ConvergedDatabase/tree/master/MTJSONIMACO>

18 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part1of3.pdf

19 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part2of3.pdf

20 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part3of3.pdf





ORACLE

# Oracle Database MultiTenant/MultiModel/InMemory/ACO Workshop

Juan Carlos Díaz

José Vázquez

Francisco Alvarez

Francisco Rivas

Madrid 19 de Enero de 2022

# Overview

- Multitenant
- MultiModel - JSON
- InMemory
- ACO

# New features brought by CDB architecture

- ✓ **Fast database provisioning** : ~1 minute vs 45 minutes with old architecture
- ✓ **Online PDB relocation** : 2 simple SQL commands. Powerful migration option.
- ✓ **Online PDB clone (local or remote)**: 1 or 2 simple SQL commands. RMAN not needed
- ✓ **PDB Snapshot clone**: 1 simple SQL command. Immediate. Space saver.
- ✓ **ASM Split Mirror**: 2 simple SQL commands. Immediate read/write clone.
- ✓ **Many B&R features** (PDB level recovery, PDB level flashback, etc)

# Overview

- MultiTenant
- **Multimodel - JSON**
- InMemory

# What is JSON?

JSON stands for **JavaScript Object Notation**

It's a "lightweight", readable data interchange format that's language independent

Most popular data format for new web applications

Instead of creating an entity relationship model to define all of the data the application needs and then mapping it to a set of relational tables

Storing JSON documents in the database greatly simplifies application development as the same schema-less data representation can be use in Application and the Database

# What is JSON?

```
{ "id": 1,
  "name": "Century 16",
  "location": { "street": "Main St",
                "city": "Redwood",
                "zipCode": "94607",
                "state": "CA",
                "phoneNumber": null
              },
  "ticketPrice": { "adultPrice": 14.95,
                  "childPrice": 9.95,
                  "seniorPrice": 9.95
                }
}
```

- A data format that consists of one or more **name value pairs** enclosed in curly brackets
- The **name** is always a string and is separated from the value by a colon
- A **value** can be a number, string, true, false null, an object or array
  - E.g. location is an **object** as it has random set of name value pairs nested inside , enclosed in { }
  - An **array** is an ordered list of related items which could be JSON objects and is enclosed in [ ]
- Each pair is separated by a comma

# Storing JSON in the Oracle Database

Table containing JSON documents

```
CREATE TABLE theater
(
  theater_id    VARCHAR2(255) ,
  json_document BLOB

);
```

- Oracle stores JSON in table columns
  - No special data type
  - Can be VARCHAR2, BLOB or CLOB
- JSON supported by all Oracle features
  - Analytics, Encryption, In-Memory, RAC, Replication, Parallel SQL, ...
  - Plus can index any JSON element

# Storing JSON in the Oracle Database

## Table containing JSON documents

```
CREATE TABLE theater
(
  theater_id    VARCHAR2(255) ,
  json_document BLOB
  CONSTRAINT is_json CHECK
    (json_document IS JSON)
);
```

```
{id:1,
 name:"Century 16",
}
{"id":1,
 "name":"Century 16",
}
```



- Oracle stores JSON in table columns
  - No special data type
  - Can be VARCHAR2, BLOB or CLOB
- JSON supported by all Oracle features
  - Analytics, Encryption, In-Memory, RAC, Replication, Parallel SQL, ...
  - Plus can index any JSON element
- **IS JSON** check constraint enforces lax JSON syntax by default
  - Does not require NAME attributes to be in double quotes



# Storing JSON in the Oracle Database

Which data type to pick?

{JSON}

## **VARCHAR2**

Best performance, easy to retrieve via SELECT

Limited max size of 32k only (with MAX\_STRING\_SIZE=EXTENDED)

## **BLOB**

Best LOB performance but not as fast as VARCHAR2

Unlimited size, not as easy to retrieve via SELECT

No potential charset conversion

## **CLOB**

Unlimited size, easy to retrieve via SELECT

Potential charset conversion (from 1 byte UTF-8 to 2 byte USC-2, double space)

Potential bigger size on disk

# Native SQL Support for JSON

## 12.2 JSON

Table containing JSON documents

```
SQL> desc THEATER
```

NAME	TYPE
JSON_DOCUMENT	BLOB

```
{ "id": 1,  
  "name": "Century 16",  
  "location": { "street": "Main St",  
                "city": "Redwood",  
                "zipCode": "94607",  
                "state": "CA",  
                "phoneNumber": null  
              }  
}
```

- JSON can be queried using simple SQL dot notation, requires **IS JSON** check constraint and table alias

```
SELECT
```

```
    t.json_document.location.city
```

```
FROM    theater t;
```

Location

-----

Los Angeles

New York

San Francisco

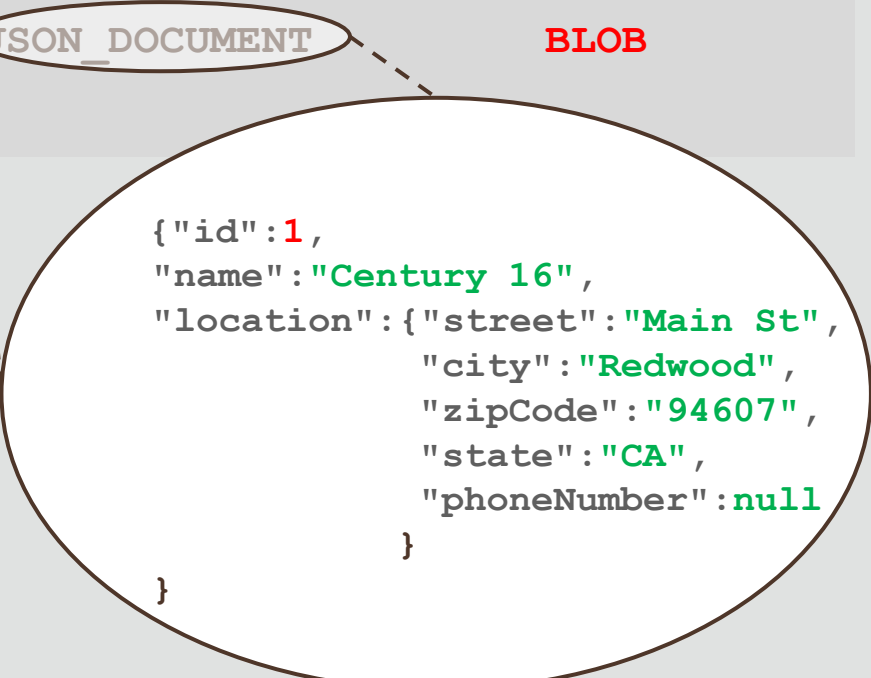
Redwood

# Alternative Mechanisms for Querying JSON

Table containing JSON documents

```
SQL> desc THEATER
```

NAME	TYPE
JSON_DOCUMENT	BLOB



```
{ "id": 1,
  "name": "Century 16",
  "location": { "street": "Main St",
                "city": "Redwood",
                "zipCode": "94607",
                "state": "CA",
                "phoneNumber": null
              }
}
```

- Without **IS JSON** constraint you need to use the **JSON\_VALUE** function to query JSON
- ```
SELECT
JSON_VALUE(t.json_document,
            '$.location.city') City
FROM theater t;
```
- City
- 
- Los Angeles  
New York  
San Francisco  
Redwood

# Oracle REST Data Service


A woman with dark hair, wearing a blue blazer and a watch, is leaning forward and listening intently to a man in a grey suit. The man is gesturing with his hands while speaking. They are in a modern office environment with large windows in the background. A laptop is open on a table in the foreground.



# Oracle Database Cloud Service

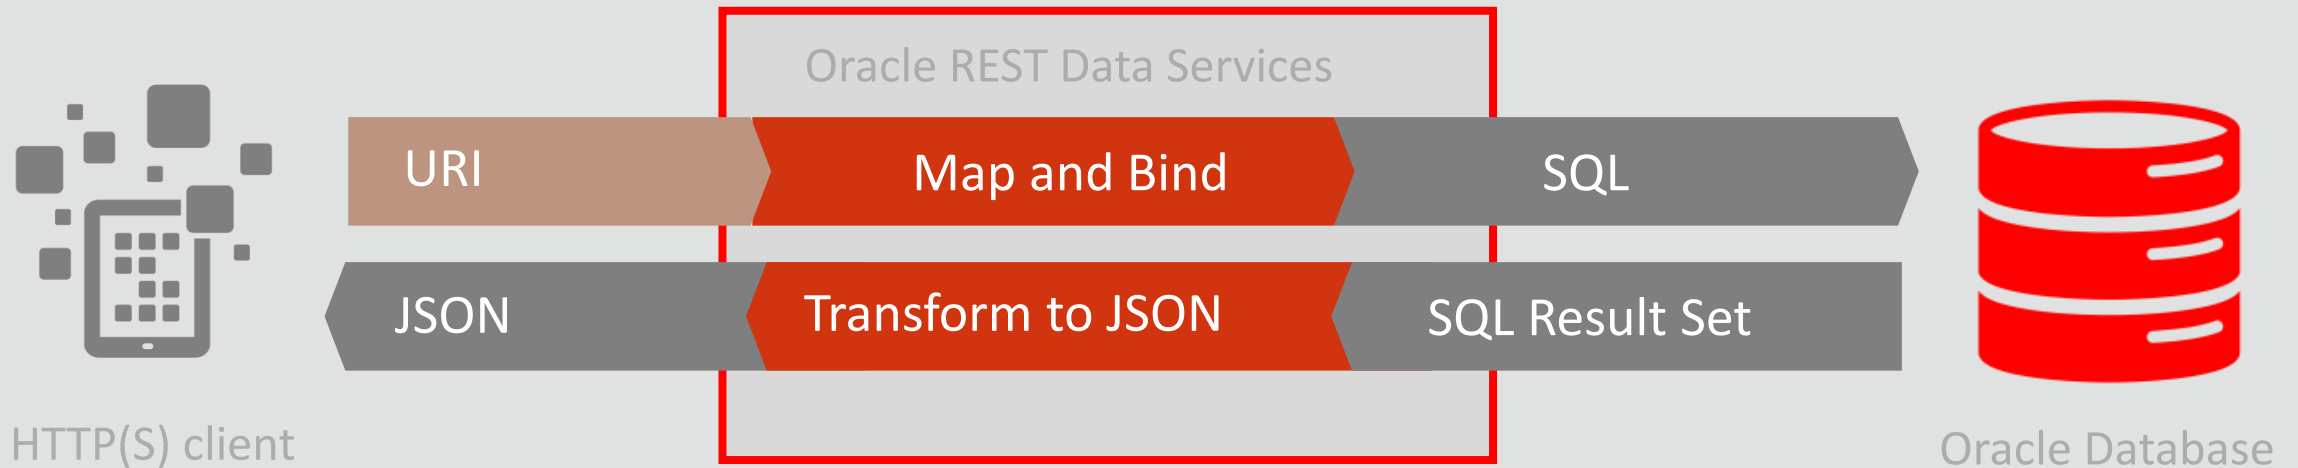
## Oracle REST Data Services

- Available **Standalone (Jetty)**, Weblogic, Tomcat & Glassfish
- Turns Database Service into an RESTFul API service
- Fully provisioned and functional in all cloud editions
- Available in 11g, 12c and 18c, **no extra cost**
- Allows publishing of URI based access to Oracle database over REST
- Results in JSON or CSV
- Mapping of URI to SQL or PL/SQL
- All HTML methods GET, PUT, POST, DELETE, PATCH
- OAuth2 integration
- Highly scalable, can use all feature of database



# Oracle REST Data Services

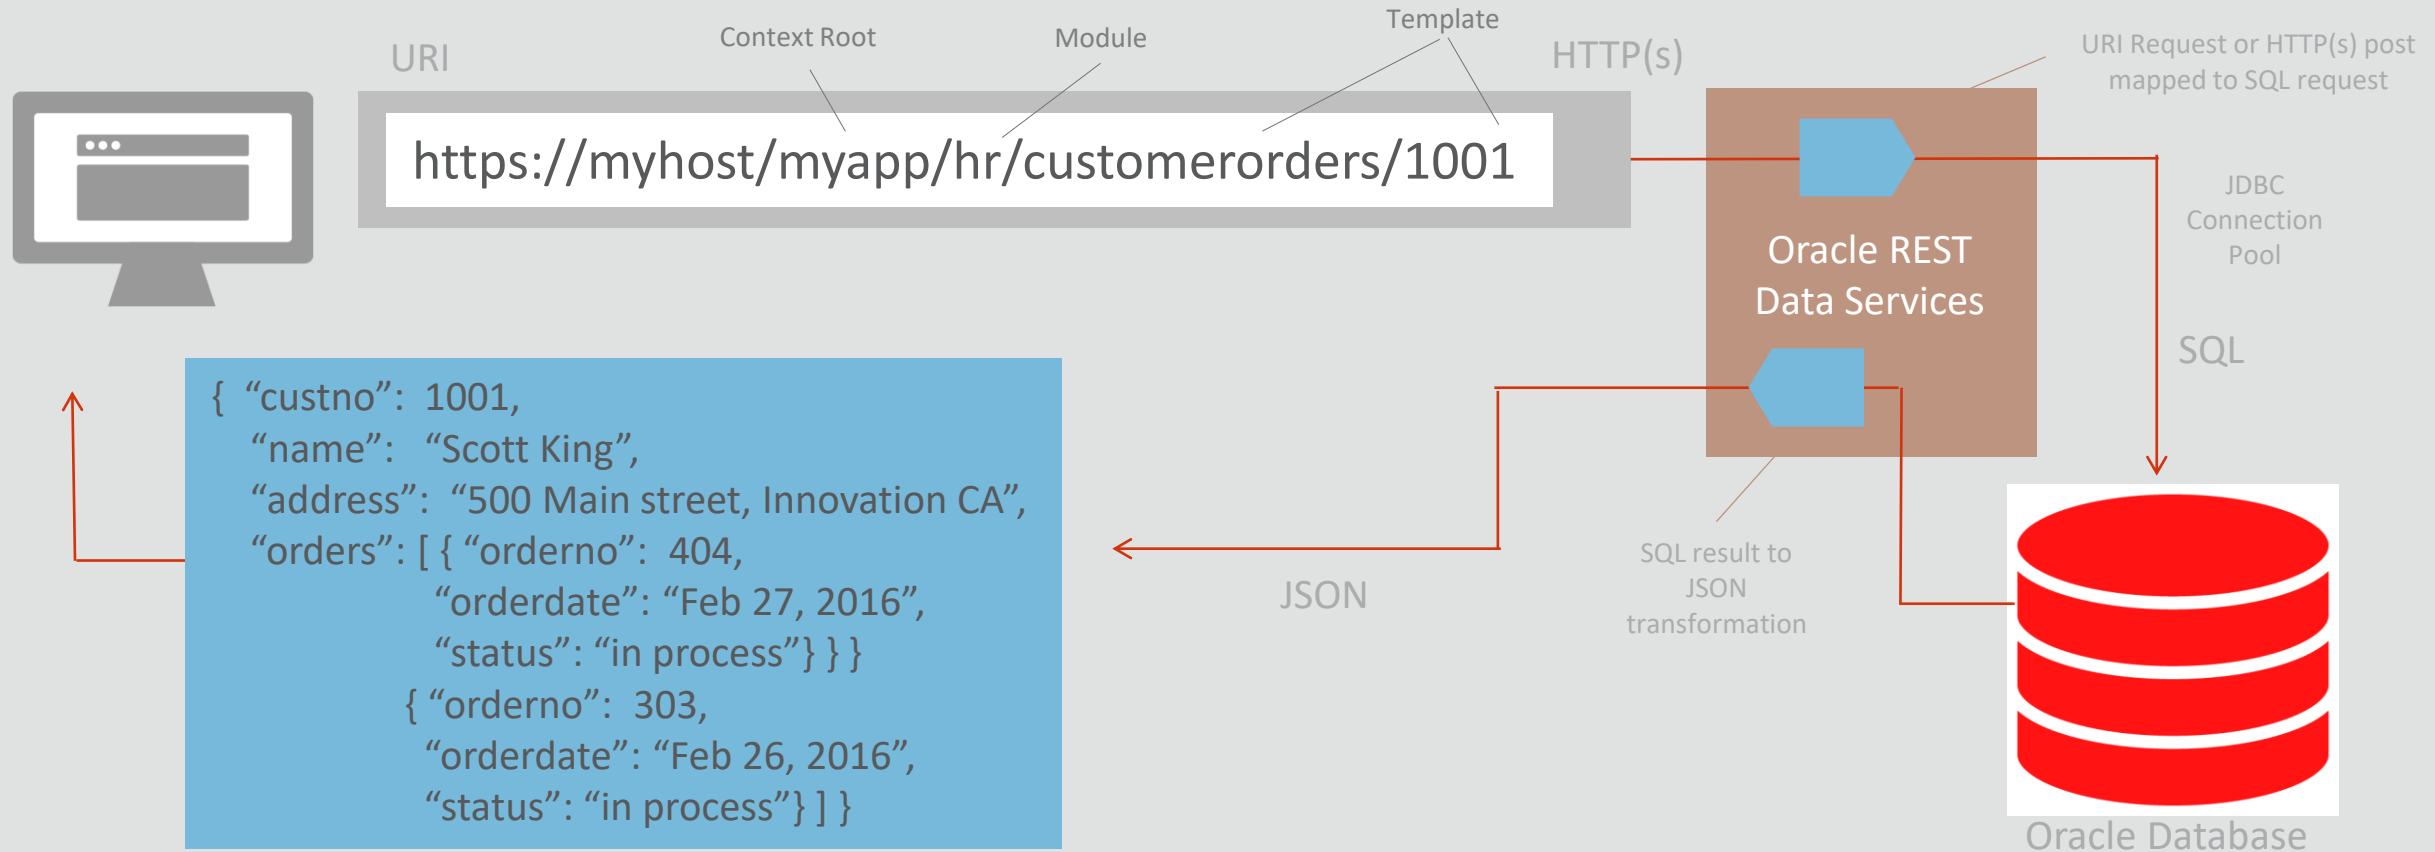
Serving JSON results from relational data



- Data stored in standard relational tables and columns
- Oracle REST Data Services (ORDS) Developer defines URI<>SQL mapping
- App Developer calls named URI over HTTP(S) gets and posts

# Oracle REST Data Services

## HTTP(s) API App-Dev with Relational Tables in Oracle Database



ORDS maps standard URI requests to corresponding relational SQL (not schemaless): e.g. SQL SELECT from customers and orders table.

ORDS also transforms the SQL results into JavaScript Object Notation (JSON), other formats include HTML, binary and CSV.

Fully committed to supporting any and all standards required by Fusion / SaaS / FMW; we are actively engaged in the ongoing dialog.

# SODA: Simple Oracle Document Access



A simple NoSQL-style API for Oracle

- Collection Management: Create and drop collections

- Document Management: CRUD (Create, Retrieve, Update and Delete) operations

- List and Search: (Query-by-Example) operations on collections

- Utility and Control: Bulk Insert, index management

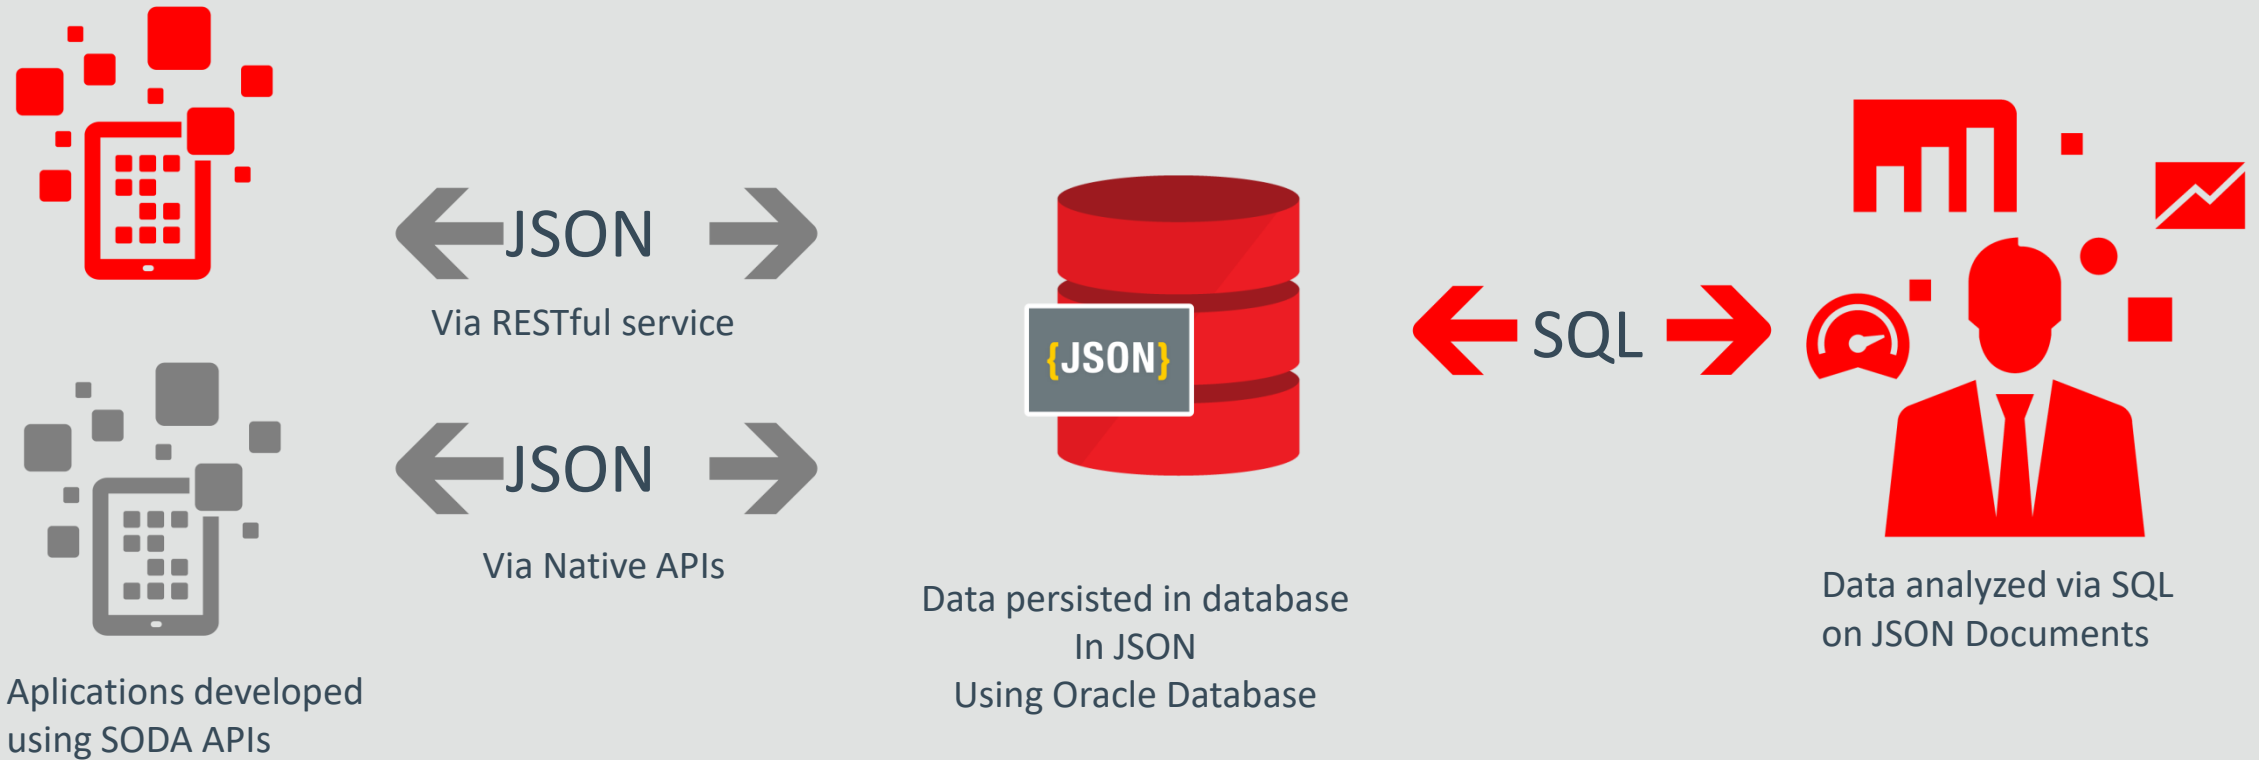
Developers can work with Oracle without learning SQL or requiring DBA support

- Same development experience as pure-play document stores



# Oracle Database as a Document Store

Flexible Schema development



# SODA for REST

```
curl --digest -u SCOTT:tiger -X POST -H "Accept: application/json" -H "Content-type: application/json" --upload-file po.json http://localhost:8080/DBJSON/SCOTT/MyCollection
```

```
{
  "items": [
    {
      "id": "A450557094D04957B36346F630CDDF9A",
      "etag": "C1354F27A5180FF7B828F01CBBC84022DCF5F7209DBF0E6DFFCC626E3B0400C3",
      "lastModified": "2015-02-09T01:03:48.291462",
      "created": "2015-02-09T01:03:48.291462"
    }
  ],
  "hasMore": false,
  "count": 1
}
```

# SODA for REST

```
curl --digest -X GET      --write-out "%{http code}\n" -u SCOTT:tiger
http://localhost:8080/DBJSON/SCOTT/Collection1/14E6656206114A12950ABF745C309CC5
```

```
{
  "PONumber": 14,
  "Reference": "SVOLLMAN-20140525",
  "Requestor": "Shanta Vollman",
  "User": "SVOLLMAN",
  "CostCenter": "A50",
  "ShippingInstructions": {
    "name": "Shanta Vollman",
    "Address": {
      "street": "200 Sporting Green",
      "city": "South San Francisco",
      "state": "CA",
      "zipCode": 99236,
      "country": "United States of America"
    },
    "Phone": [
      {
        "type": "Office",
        "number": "823-555-9969"
      }
    ]
  },
  "Special Instructions": "Counter to Counter",
  "LineItems": [...]
}
```

# Workshop

## MultiTenant (Advanced) + JSON

<https://github.com/OracleDataManagementSpain/ConvergedDatabase/tree/master/MTJSONIMACO>

18 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part1of3.pdf

19 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part2of3.pdf

20 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part3of3.pdf



ORACLE

# Oracle Database MultiTenant/MultiModel/InMemory/ACO Workshop

Juan Carlos Díaz

José Vázquez

Francisco Alvarez

Francisco Rivas

Madrid 20 de Enero de 2022

# Overview

- Multitenant
- MultiModel - JSON
- Oracle In-Memory Database
- ACO

# Overview

- MultiTenant
- MultiModel - JSON
- Oracle Database In-Memory



# Oracle Database In-Memory Goals

Real Time  
Analytics



100x

Accelerate Mixed  
Workload OLTP




2x

No Changes to  
Applications



Trivial to  
Implement



# Row Format Databases vs. Column Format Databases

Row



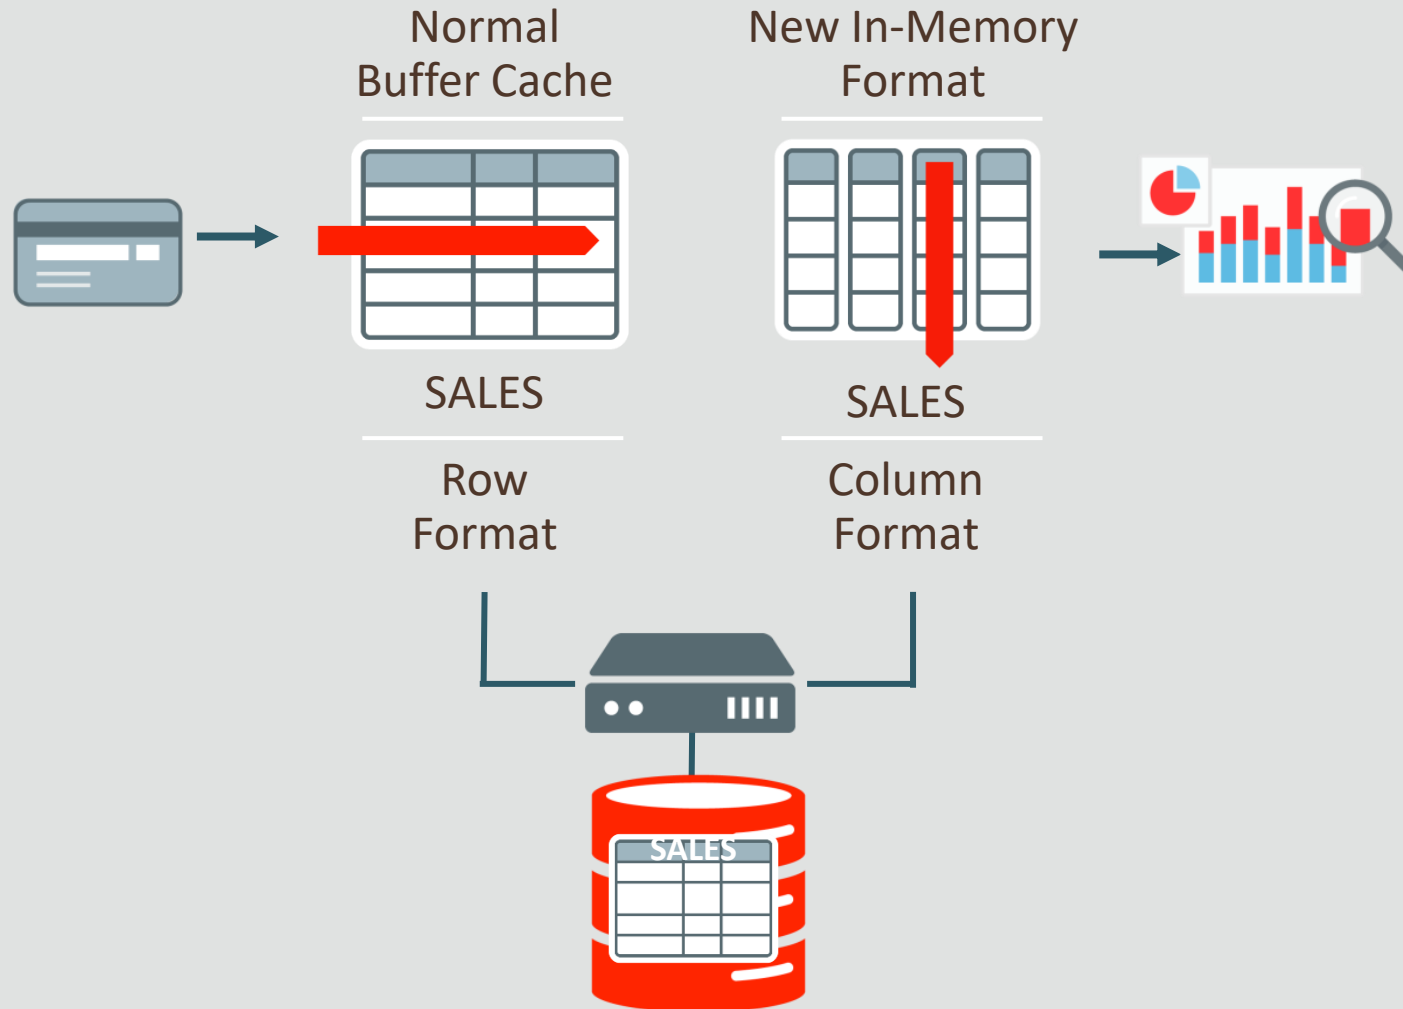
- **Transactions** run faster on row format
  - Example: Query or Insert a sales order
  - Fast processing few rows, many columns

Column



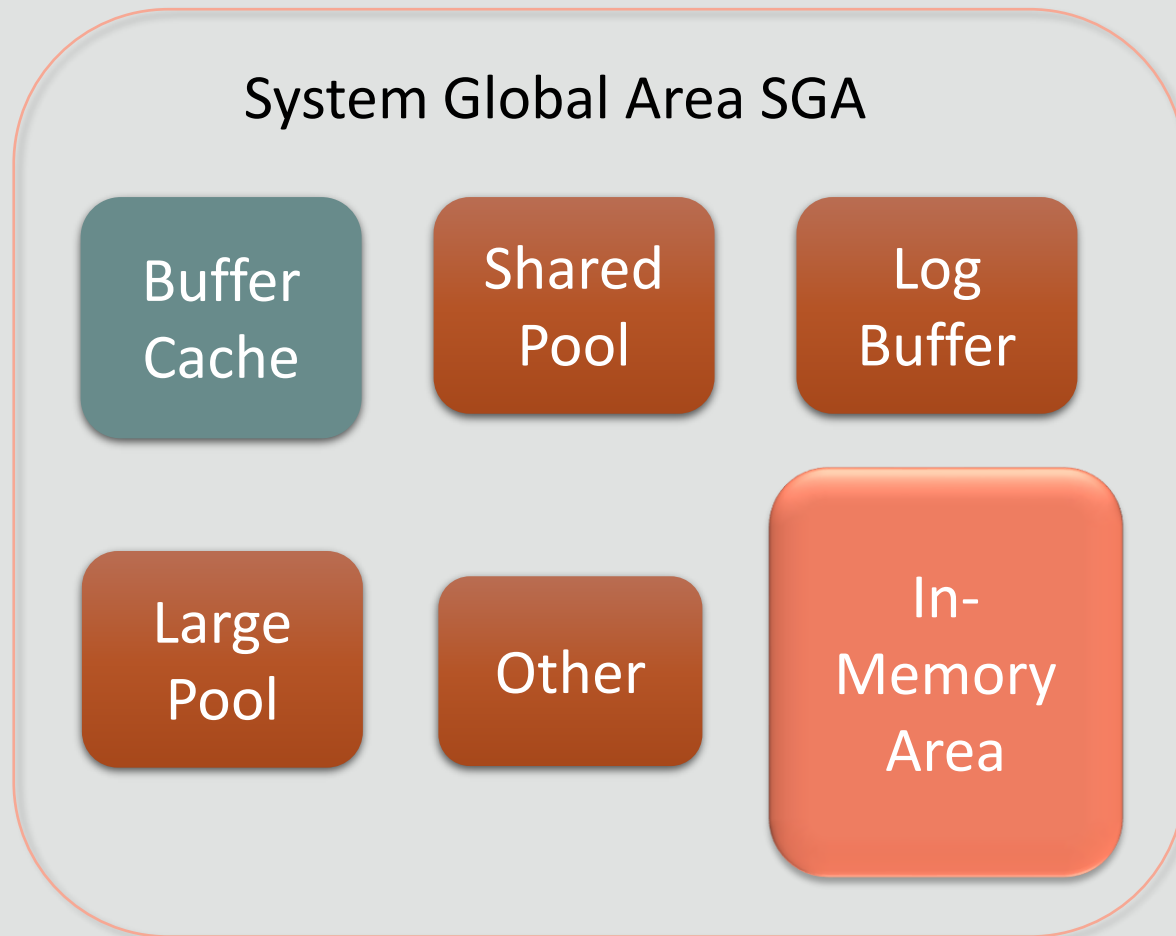
- **Analytics** run faster on column format
  - Example : Report on sales totals by region
  - Fast accessing few columns, many rows

# Breakthrough: Dual Format Database



- **BOTH** row and column formats for same table
- Simultaneously active and transactionally consistent
- Analytics & reporting use new in-memory Column format
- OLTP uses proven row format

# In-Memory Area: Static Area within SGA



Contains data in the new In-Memory Columnar Format

Controlled by INMEMORY\_SIZE parameter

- Minimum size of 100MB

Can be re-sized larger while database is running (12.2)

SGA\_TARGET must be large enough to accommodate In-Memory area

# Oracle In-Memory: Simple to Implement

1. Configure Memory Capacity

```
inmemory_size = XXX GB
```

2. Configure tables or partitions to be in memory

```
alter table | partition ... inmemory;
```

3. Later drop analytic indexes to speed up OLTP

# Populating : In-Memory Column Store

```
ALTER TABLE sales INMEMORY;
```

```
ALTER TABLE sales NO INMEMORY;
```

```
CREATE TABLE customers .....  
PARTITION BY LIST  
  (PARTITION p1 ..... INMEMORY,  
   (PARTITION p2 ..... NO INMEMORY) ;
```

- New INMEMORY ATTRIBUTE
- Following segment types are eligible
  - Tables
  - Partitions
  - Subpartition
  - Materialized views
- Following segment types not eligible
  - IOTs
  - Hash clusters
  - Out of line LOBs

Pure OLTP  
Features

# Populating : In-Memory Column Store

```
ALTER TABLE sales INMEMORY  
NO INMEMORY (PROD_ID) ;
```

```
CREATE TABLE orders  
  (c1 number,  
   c2 varchar(20) ,  
   c3 number)  
INMEMORY PRIORITY CRITICAL  
NO INMEMORY (c1) ;
```

- Possible to populate only certain columns from a table or partition
- Order in which objects are populated controlled by PRIORITY subclause
  - Critical, high, medium, low
  - Default – none (populate on first access)
  - Does not control the speed of population

# Populating : In-Memory Column Store

```
ALTER MATERIALIZED VIEW mv1  
INMEMORY MEMCOMPRESS FOR QUERY;
```

```
CREATE TABLE trades  
  (Name varchar(20) ,  
   Desc varchar(200) )  
INMEMORY  
MEMCOMPRESS FOR DML(desc) ;
```

- Objects compressed during population
- New compression techniques
  - Focused on scan performance
- Controlled by MEMCOMPRESS subclause
- Multiple levels of compression
- Possible to use a different level for different partitions in a table



# Populating : In-Memory Column Store

```
CREATE TABLE ORDERS .....  
PARTITION BY RANGE .....  
  (PARTITION p1 .....  
    INMEMORY NO MEMCOMPRESS  
  PARTITION p2 .....  
    INMEMORY MEMCOMPRESS FOR DML,  
  PARTITION p3 .....  
    INMEMORY MEMCOMPRESS FOR QUERY,  
  :  
  PARTITION p200 .....  
    INMEMORY MEMCOMPRESS FOR CAPACITY  
  ) ;
```

- Different levels
  - FOR DML  
Use on tables or partitions with very active DML activity
  - FOR QUERY  
Default mode for most tables
  - FOR CAPACITY  
For less frequently accessed segments
- Easy to switch levels as part of ILM strategy

# Oracle Compression Advisor And In-Memory

```
DECLARE
  l_blkcnt_cmp      BINARY_INTEGER;
  l_blkcnt_uncmp    BINARY_INTEGER;
  l_row_cmp         BINARY_INTEGER;
  l_row_uncmp       BINARY_INTEGER;
  l_cmp_ratio       NUMBER;
  l_comptype_str    VARCHAR2(100);
BEGIN
  dbms_compression.get_compression_ratio(
    -- input parameters
    scratchtbsname => 'USERS',           -- scratch tablespace
    ownname        => 'SSB',             -- owner of the table
    objname        => 'LINEORDER',       -- table name
    subobjname     => NULL,              -- partition name
    comptype       => DBMS_COMPRESSION.COMP_INMEMORY_QUERY, -- compression algorithm
    -- output parameters
    blkcnt_cmp     => l_blkcnt_cmp,      -- number of compressed blocks
    blkcnt_uncmp   => l_blkcnt_uncmp,    -- number of uncompressed blocks
    row_cmp        => l_row_cmp,         -- number of rows in a compressed block
    row_uncmp      => l_row_uncmp,       -- number of rows in an uncompressed block
    cmp_ratio      => l_cmp_ratio,       -- compression ratio
    comptype_str   => l_comptype_str     -- compression type
  );
  dbms_output.put_line('LINEORDER '||l_comptype_str||' ratio: '||to_char(l_cmp_ratio,'99.999'));
END;
```

- Easy way to determine memory requirements
- Use DBMS\_COMPRESSION
- Applies MEMCOMPRESS to sample set of data from a table
- Returns estimated compression ratio


# Oracle In-Memory Advisor

| Object Type | Object                     | Estimated In-Memory Size | Analytics Processing Seconds | Estimated Reduced Analytics Processing Seconds | Estimated Analytics Processing Performance Improvement Factor | Benefit / Cost Ratio (Improvement Factor / In-Memory Size) |
|-------------|----------------------------|--------------------------|------------------------------|------------------------------------------------|---------------------------------------------------------------|------------------------------------------------------------|
| Table       | SOE.LOGON                  | 451.76MB                 | 2114                         | 1,887                                          | 9.3X                                                          | 20.586                                                     |
| Table       | SOE.CARD_DETAILS           | 607.32MB                 | 8346                         | 7,248                                          | 7.6X                                                          | 12.514                                                     |
| Table       | SOE.ADDRESSES              | 1.09GB                   | 5237                         | 4,621                                          | 8.5X                                                          | 7.798                                                      |
| Partition   | SOE.PRODUCT MOCKUP.Y2014Q1 | 812.6MB                  | 2003                         | 1,489                                          | 3.9X                                                          | 4.799                                                      |
| Table       | SOE.CUSTOMERS              | 1.10GB                   | 108                          | 95                                             | 8.2X                                                          | 7.455                                                      |
| Table       | SOE.ORDER_ITEMS            | 2.19GB                   | 7128                         | 6,393                                          | 9.7X                                                          | 4.429                                                      |
| Table       | SOE.ORDERS                 | 1.34GB                   | 3512                         | 2,917                                          | 5.9X                                                          | 4.403                                                      |
| Table       | SOE.PRODUCT_INFORMATION    | 1.78MB                   | 2873                         | 2,205                                          | 4.3X                                                          | 2.416                                                      |
| Partition   | SOE.PRODUCT MOCKUP.Y2013Q4 | 1.62GB                   | 97                           | 1,489                                          | 3.7X                                                          | 2.284                                                      |
| Partition   | SOE.PRODUCT MOCKUP.Y2014Q2 | 3.37GB                   | 642                          | 493                                            | 4.3X                                                          | 1.276                                                      |

- New In-Memory Advisor
- Analyzes existing DB workload via AWR & ASH repositories
- Provides list of objects that would benefit most from being populated into IM column store

# Why is an In-Memory scan faster than the buffer cache?

Buffer Cache



Row Format

```
SELECT COL4 FROM MYTABLE;
```



RESULT

# Why is an In-Memory scan faster than the buffer cache?

IM Column Store

| COL1 | COL2 | COL3 | COL4 |
|------|------|------|------|
| X    | X    | X    | X    |
| X    | X    | X    | X    |
| X    | X    | X    | X    |
| X    | X    | X    | X    |
| X    | X    | X    | X    |

Column Format

X  
X  
X  
X  
X


SELECT **COL4** FROM MYTABLE;



RESULT

# Oracle In-Memory Column Store Storage Index

Example: Find sales from stores with a store\_id of 8 or higher



- Each column is made up of multiple column units
- Min / max value is recorded for each column unit in a storage index
- Storage index provides partition pruning like performance for **ALL** queries

# SQL statement is scanning data in the IM ?

```
SELECT Max(lo_ordtotalprice) most_expensive_order  
FROM lineorder  
WHERE lo_shipmode = 5;
```

| Id | Operation                   | Name      |
|----|-----------------------------|-----------|
| 0  | SELECT STATEMENT            |           |
| 1  | SORT AGGREGATE              |           |
| 2  | TABLE ACCESS IN MEMORY FULL | LINEORDER |



# Oracle In-Memory Requires Zero Application Changes

**Full Functionality**

**Easy to Implement**

**Fully Compatible**

**Fully Multitenant**

- **ZERO restrictions** on SQL
- No migration of data
- All existing applications run unchanged
- Oracle In-Memory is Cloud Ready

**ORACLE®**  
E-BUSINESS SUITE

**ORACLE®**  
FUSION APPLICATIONS

**ORACLE®**  
JD EDWARDS

**ORACLE®**  
PEOPLESOFT

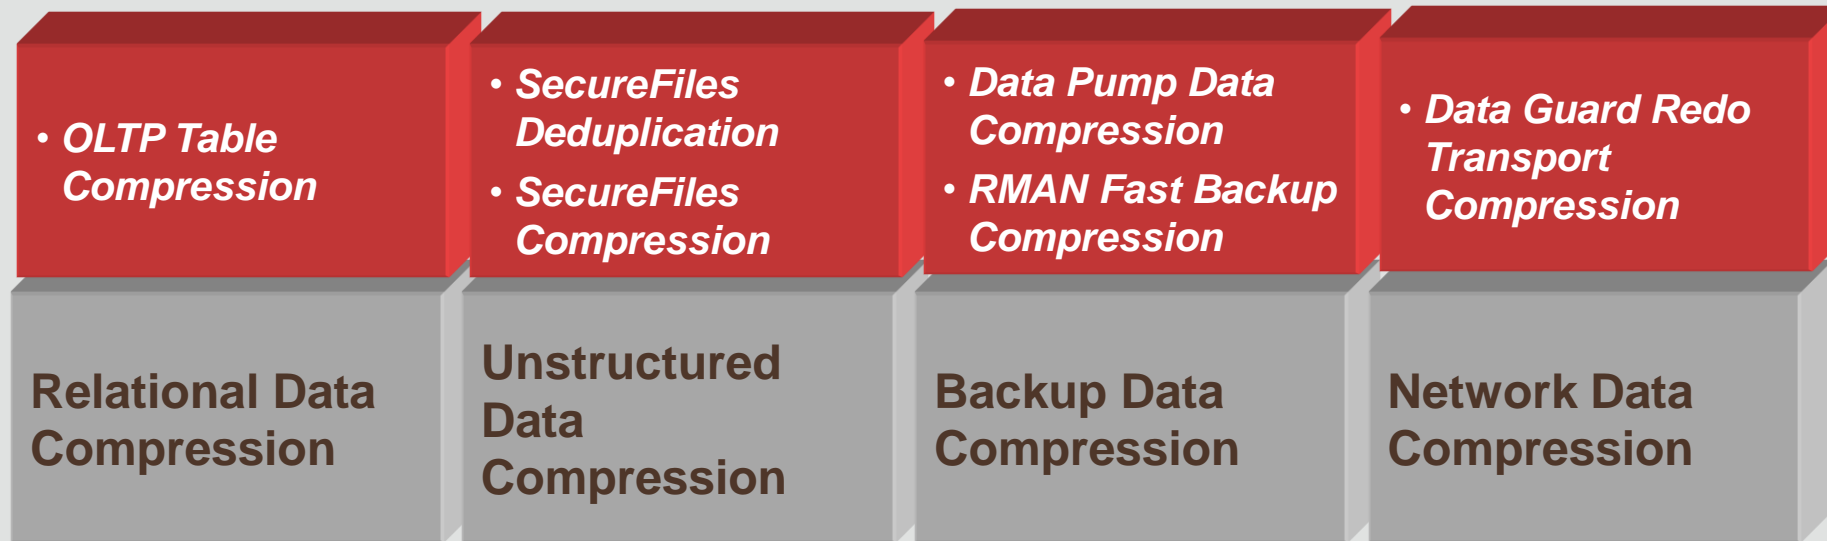
**ORACLE®**  
SIEBEL

**Uniquely Achieves All In-Memory Benefits With No Application Changes**

# Overview

- MultiTenant
- JSON
- InMemory
- ACO

# Oracle Advanced Compression Option



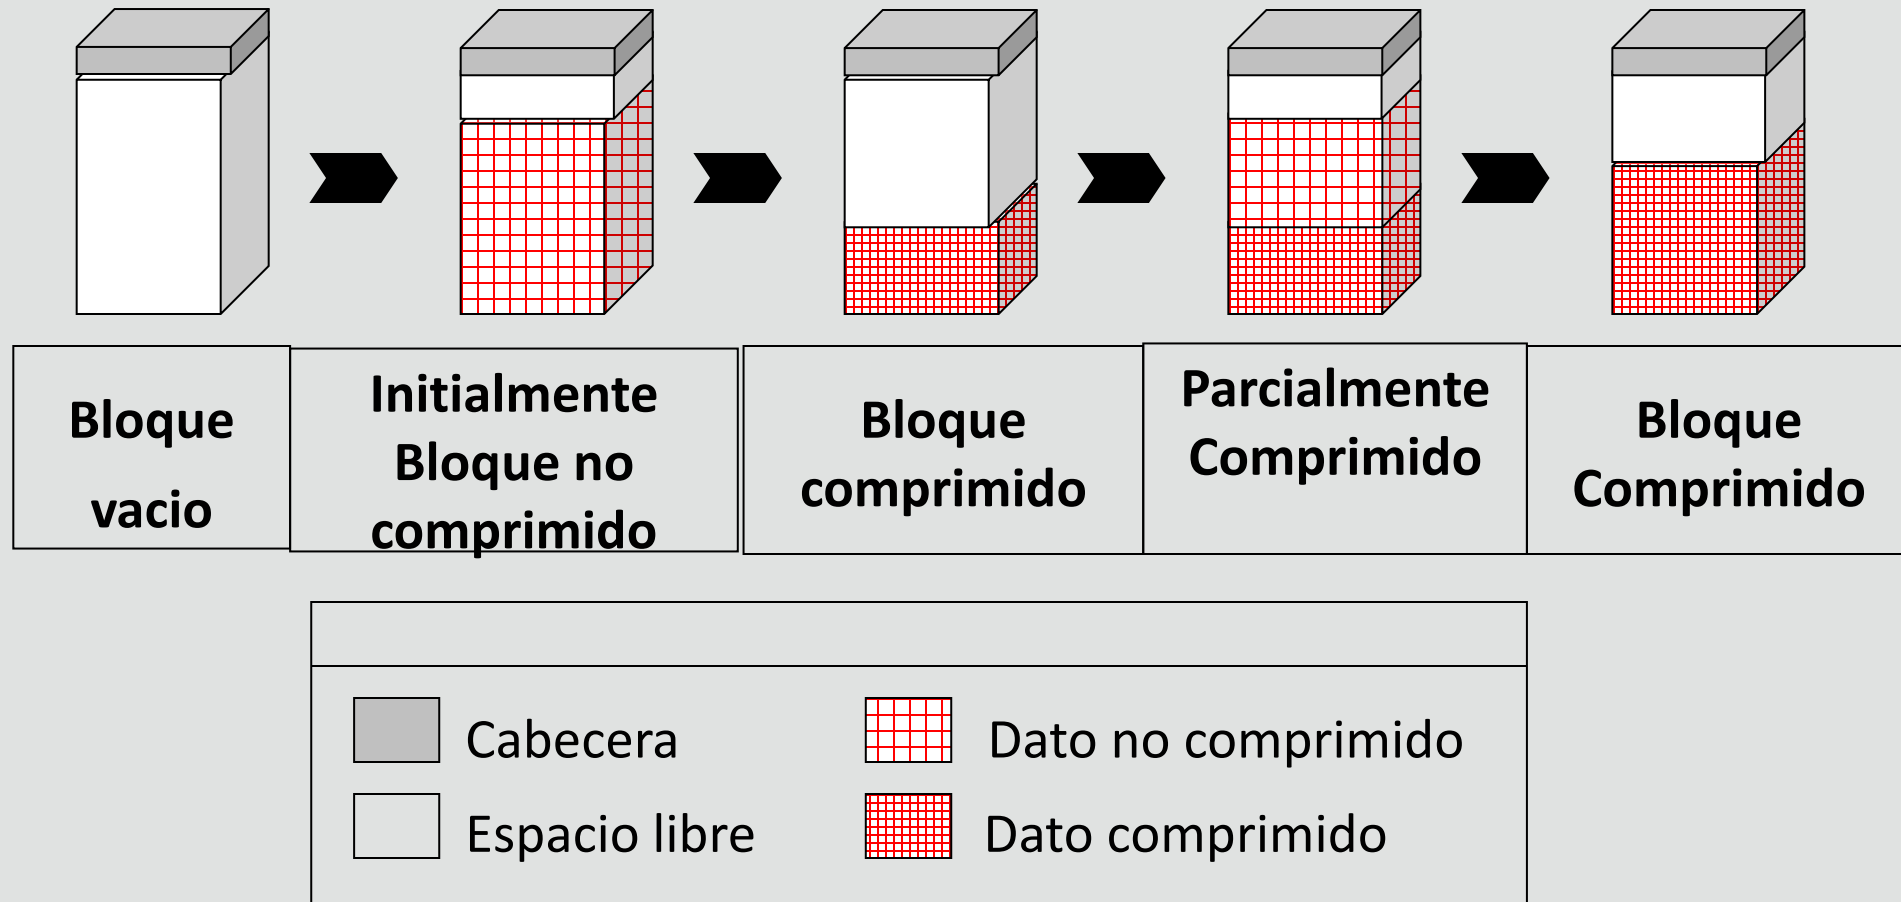
Reduce requerimientos de recursos y costes

Almacenamiento

Ancho de banda

Uso de memoria

# OLTP Proceso compresión de tabla



# OLTP Compresión Tabla

## Employee Table

| ID | FIRST_NAME | LAST_NAME |
|----|------------|-----------|
| 1  | John       | Doe       |
| 2  | Jane       | Doe       |
| 3  | John       | Smith     |
| 4  | Jane       | Doe       |

## Initially Uncompressed Block

| Header       |            |  |
|--------------|------------|--|
| 1•John•Doe   | 2•Jane•Doe |  |
| 3•John•Smith | 4•Jane•Doe |  |

Free Space

```
INSERT INTO EMPLOYEE
VALUES (5, 'Jack', 'Smith');
COMMIT;
```

# OLTP Compresión Tabla

**Employee Table**

| ID | FIRST_NAME | LAST_NAME |
|----|------------|-----------|
| 1  | John       | Doe       |
| 2  | Jane       | Doe       |
| 3  | John       | Smith     |
| 4  | Jane       | Doe       |
| 5  | Jack       | Smith     |

**Local  
Symbol Table**

**Compressed Block  
Header**

John=**0** | Doe=**1** | Jane=**2** | Smith=**3**

1 • **0** • **1** 2 • **2** • **1** 3 • **0** • **3** 4 •

2 • **1** 5 • Jack • **3**

Free Space

# OLTP Table Compression

## Uncompressed Block Header

# Header

|            |               |    |  |
|------------|---------------|----|--|
| 1•John•Doe | 2•Jane•       |    |  |
| Doe        | 3•John•Smith  | 4• |  |
| Jane • Doe | 5•Jack •Smith |    |  |
| Free Space |               |    |  |

Local  
Symbol Table

## Compressed Block Header

|                                   |       |          |    |
|-----------------------------------|-------|----------|----|
| Compressed Block Header           |       |          |    |
| John=0   Doe=1   Jane=2   Smith=3 |       |          |    |
| 1•0•1                             | 2•2•1 | 3•0•3    | 4• |
| 2•                                | 1     | 5•Jack•3 |    |
| Free Space                        |       |          |    |


More Data  
Per Block



# 19c Oracle SecureFiles

*Gestión de datos consolidada y segura*

- **SecureFiles** información no estructurada o semi estructurada almacenada en la BBDD.
- **Reingeniería de los LOBs**. Acceso más rápido y con mas funcionalidad.
  - Eliminación de elementos duplicados de manera transparente (deduplication), **compresión** y **encriptación**.
  - Utiliza las funcionalidades de **seguridad**, **disponibilidad**,
  - y **escalabilidad** de la base de datos.
  - Migración sencilla de LOBs.



# Compresión Data Pump

Oracle Database 12c extiende la compresión de la tabla durante los exports

No necesidad de descomprimir antes de importar

```
COMPRESSION={ALL | DATA_ONLY |  
[METADATA_ONLY] | NONE}
```

Simple compresión de ambos data y metadata

Datos comprimidos directamente a disco resultando una disminución de espacio

# Compresión Backup

- RMAN Backup Compression

- Backup es comprimido antes de escribirse a disco
- No necesario descomprimir antes de restauración
- Compresion nivel LOW Algoritmo de compresión más rápido
  - Más adecuado cuando hay limites en CPU
- Nivel de compresión MEDIUM
  - Balancea entre uso de CPU y ratio de compresión
- Compresión LEVEL HIGH Mejor ratio de compresión y más alto utilización de CPU
  - Más adecuado cuando hay límites de red o I/O

# Compresión de Red

## Data Guard Redo Transport Services

- Data Guard Redo Transport Services transfiere datos de log a base standby
- Con Advanced Compression, redo data pueden ser transmitidos en formato comprimido para reducir consumo de ancho banda y tiempo de transmisión
- 12c synchronous redo transport (SYNC) o asynchronous redo transport (ASYNC).
- Ejemplo:
  - `LOG_ARCHIVE_DEST_3='SERVICE=denver SYNC COMPRESSION=ENABLE | [DISABLE] '`

# Workshop

## IN-Memory Database + Advanced Compression Option

<https://github.com/OracleDataManagementSpain/ConvergedDatabase/tree/master/MTJSONIMACO>

18 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part1of3.pdf

19 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part2of3.pdf

20 Enero

WORKSHOP\_Multitenant\_Multimodel\_InMemory\_HOL\_v3\_part3of3.pdf

ORACLE

# ENCUESTA

*Workshop Virtual BD Convergente:  
Multitenant, Multimodel, In-memory*



¡Tu opinión es muy importante!

Escanea el QR para responder o entra aquí:

<https://bit.ly/3AbzRot>



# Thank you

---

**Oracle Spain**

